

# CENG403 - Computer Vision Quiz: Homework set THE-2

## Coding Questions - Set 2

Student Name: \_\_\_\_\_

### Problem 9

#### Deformable Convolution Core Implementation

Complete the deformable convolution forward pass structure:

```
1 def deform_conv2d_np(input_map, offset, mask, weight, stride=1, padding=0):
2     """
3     input_map: (N, C_in, H_in, W_in)
4     offset: (N, 2*K*K, H_out, W_out) - contains dy, dx pairs
5     mask: (N, K*K, H_out, W_out) - modulation scalars
6     weight: (C_out, C_in, K, K)
7     """
8     N, C_in, H_in, W_in = input_map.shape
9     C_out, _, K, _ = weight.shape
10
11     # Calculate output dimensions
12     H_out = (H_in + 2*padding - K) // stride + 1
13     W_out = (W_in + 2*padding - K) // stride + 1
14
15     # Pad input if needed
16     if padding > 0:
17         padded_input = -----
18     else:
19         padded_input = input_map
20
21     # Initialize output
22     output = np.zeros((N, C_out, H_out, W_out))
23
24     # Main computation loops
25     for n in range(N):
26         for c_out in range(C_out):
27             for h_out in range(H_out):
28                 for w_out in range(W_out):
29                     for k_h in range(K):
30                         for k_w in range(K):
31                             k_idx = k_h * K + k_w
32
33                             # Get offset for this kernel position
34                             dy = offset[n, 2*k_idx, h_out, w_out]
35                             dx = offset[n, 2*k_idx+1, h_out, w_out]
36
37                             # Get modulation mask
38                             m_k = mask[n, k_idx, h_out, w_out]
39
40                             # Calculate sampling position
41                             sample_y = h_out * stride + k_h + dy
42                             sample_x = w_out * stride + k_w + dx
43
44                             # For each input channel
45                             for c_in in range(C_in):
46                                 # Bilinear interpolation
47                                 interp_val = bilinear_interpolate(
48                                     padded_input[n, c_in], sample_y, sample_x
49                                 )
50
51                                 # Apply convolution weight and modulation
```

```
52         output[n, c_out, h_out, w_out] += -----
53
54     return output
```

Complete the final line: -----

## Problem 10

### Gradient Computation with torch.autograd

Complete the manual gradient computation:

```
1 # RNN forward pass
2 h = torch.zeros(H, requires_grad=True)
3 logits_list = []
4
5 for t in range(seq_len):
6     h = torch.tanh(W_xh @ inputs[t] + W_hh @ h + b_xh + b_hh)
7     logits = W_hy @ h + b_y
8     logits_list.append(logits)
9
10 # Compute loss
11 all_logits = torch.stack(logits_list)
12 log_probs = F.log_softmax(all_logits, dim=1)
13 loss = F.nll_loss(log_probs, targets)
14
15 # Manual gradient computation
16 grad_W_xh = torch.autograd.grad(loss, W_xh, _____)[0]
17 grad_W_hh = torch.autograd.grad(loss, W_hh, _____)[0]
18 grad_b_xh = torch.autograd.grad(loss, b_xh, _____)[0]
19 grad_b_hh = torch.autograd.grad(loss, b_hh, _____)[0]
20 grad_W_hy = torch.autograd.grad(loss, W_hy, _____)[0]
21 grad_b_y = torch.autograd.grad(loss, b_y, _____)[0]
22
23 print("Gradients computed successfully!")
```

Fill in the missing parameter: \_\_\_\_\_

Debug this gradient computation error:

```
1 # This code will throw an error
2 for param in [W_xh, W_hh, b_xh, b_hh, W_hy, b_y]:
3     grad = torch.autograd.grad(loss, param)
4     print(f"Gradient computed: {grad.shape}")
```

Error type: \_\_\_\_\_

Why it happens: \_\_\_\_\_

How to fix: \_\_\_\_\_

## Problem 11

### BatchNorm Implementation and Usage

Complete the CNN with BatchNorm:

```
1 class CNNWithBatchNorm(nn.Module):
2     def __init__(self, num_classes=100):
3         super(CNNWithBatchNorm, self).__init__()
4
5         self.conv1 = nn.Conv2d(3, 64, 3, padding=1)
6         self.bn1 = nn.BatchNorm2d(_____)
7
8         self.conv2 = nn.Conv2d(64, 128, 3, padding=1)
9         self.bn2 = _____
10
11        self.conv3 = nn.Conv2d(128, 256, 3, padding=1)
12        self.bn3 = nn.BatchNorm2d(256)
13
14        self.pool = nn.MaxPool2d(2, 2)
15
16        self.fc1 = nn.Linear(256 * 4 * 4, 512)
17        self.fc2 = nn.Linear(512, num_classes)
18
19    def forward(self, x):
20        # Block 1
21        x = self.conv1(x)
22        x = self.bn1(x)
23        x = F.relu(x)
24        x = self.pool(x)
25
26        # Block 2
27        x = _____
28        x = _____
29        x = _____
30        x = self.pool(x)
31
32        # Block 3
33        x = self.conv3(x)
34        x = self.bn3(x)
35        x = F.relu(x)
36        x = self.pool(x)
37
38        # Classifier
39        x = x.view(x.size(0), -1)
40        x = F.relu(self.fc1(x))
41        x = self.fc2(x)
42
43        return x
```

What's wrong with this BatchNorm usage?

```
1 # During training
2 model.train()
3 for batch in train_loader:
4     outputs = model(inputs)
5     # ... training code
6
7 # During validation
8 model.eval()
9 for batch in val_loader:
```

```
10 |         with torch.no_grad():
11 |             outputs = model(inputs) # BatchNorm behaves differently here
```

**Explain the behavior difference:** .....

## Problem 12

### Training Loop with Learning Rate Scheduling

Complete the training loop with scheduler:

```
1 def train_model(model, train_loader, val_loader, num_epochs=20):
2     criterion = nn.CrossEntropyLoss()
3     optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
4     scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=7, gamma
5         =0.1)
6
7     train_losses = []
8     val_accuracies = []
9
10    for epoch in range(num_epochs):
11        # Training phase
12        model.train()
13        running_loss = 0.0
14
15        for images, labels in train_loader:
16            images, labels = images.to(device), labels.to(device)
17
18            ----- # Clear gradients
19            outputs = model(images)
20            loss = criterion(outputs, labels)
21            ----- # Backward pass
22            ----- # Update weights
23
24            running_loss += loss.item()
25
26        # Validation phase
27        model.-----
28        val_correct = 0
29        val_total = 0
30
31        with -----:
32            for images, labels in val_loader:
33                images, labels = images.to(device), labels.to(device)
34                outputs = model(images)
35
36                _, predicted = torch.max(outputs.data, 1)
37                val_total += labels.size(0)
38                val_correct += (predicted == labels).sum().item()
39
40        # Update learning rate
41        -----
42
43        # Record metrics
44        avg_loss = running_loss / len(train_loader)
45        val_acc = 100 * val_correct / val_total
46
47        train_losses.append(avg_loss)
48        val_accuracies.append(val_acc)
49
50        print(f'Epoch [{epoch+1}/{num_epochs}], Loss: {avg_loss:.4f}, Val Acc: {val_acc:.2f}%')
51
52    return train_losses, val_accuracies
```

## Problem 13

### Data Augmentation and Preprocessing

Complete the data augmentation pipeline:

```
1 import torchvision.transforms as transforms
2
3 # Training augmentations
4 train_transform = transforms.Compose([
5     transforms.RandomHorizontalFlip(p=_____),
6     transforms.RandomRotation(degrees=_____),
7     transforms.ColorJitter(brightness=0.2, contrast=0.2),
8     transforms.RandomCrop(32, padding=4),
9     transforms.ToTensor(),
10    transforms.Normalize(
11        mean=[0.5071, 0.4867, 0.4408], # CIFAR100 stats
12        std=[0.2675, 0.2565, 0.2761]
13    )
14 ])
15
16 # Validation/Test (no augmentation)
17 test_transform = transforms.Compose([
18     _____, # Convert to tensor
19     _____ # Apply same normalization
20 ])
```

Fix this data loading issue:

```
1 # Problem: validation set uses training augmentations
2 full_dataset = CIFAR100(root='./data', train=True, transform=train_transform)
3 train_set, val_set = random_split(full_dataset, [40000, 10000])
4
5 # Both train_set and val_set use train_transform!
6 # How to fix this?
```

**Solution:**

## Problem 14

### Model Evaluation and Metrics

Complete the evaluation function:

```
1 def evaluate_model(model, test_loader, device):
2     model.eval()
3
4     total_samples = 0
5     correct_top1 = 0
6     correct_top5 = 0
7
8     class_correct = list(0. for i in range(100)) # CIFAR100 has 100 classes
9     class_total = list(0. for i in range(100))
10
11     with torch.no_grad():
12         for images, labels in test_loader:
13             images, labels = images.to(device), labels.to(device)
14             outputs = model(images)
15
16             # Top-1 accuracy
17             _, pred_top1 = torch.max(outputs, 1)
18             correct_top1 += -----
19
20             # Top-5 accuracy
21             _, pred_top5 = torch.topk(outputs, 5, dim=1)
22             for i in range(labels.size(0)):
23                 if labels[i] in pred_top5[i]:
24                     correct_top5 += 1
25
26             total_samples += labels.size(0)
27
28             # Per-class accuracy
29             c = (pred_top1 == labels).squeeze()
30             for i in range(labels.size(0)):
31                 label = labels[i]
32                 class_correct[label] += c[i].item()
33                 class_total[label] += 1
34
35     # Calculate accuracies
36     top1_acc = 100 * correct_top1 / total_samples
37     top5_acc = 100 * correct_top5 / total_samples
38
39     print(f'Top-1 Accuracy: {top1_acc:.2f}%')
40     print(f'Top-5 Accuracy: {top5_acc:.2f}%')
41
42     # Per-class accuracy
43     for i in range(10): # Print first 10 classes
44         if class_total[i] > 0:
45             acc = 100 * class_correct[i] / class_total[i]
46             print(f'Class {i}: {acc:.2f}%')
47
48     return top1_acc, top5_acc
```

Fill in the top-1 accuracy calculation: -----



## Problem 15

### Debugging Common Issues

Identify and fix the bugs in these code snippets:

#### Bug 1 - Memory Issue:

```
1 # This causes memory leak during training
2 for epoch in range(100):
3     for batch_idx, (data, target) in enumerate(train_loader):
4         output = model(data)
5         loss = criterion(output, target)
6         loss.backward()
7         optimizer.step()
8
9         losses.append(loss) # Bug here!
```

Problem: \_\_\_\_\_

Fix: \_\_\_\_\_

#### Bug 2 - Device Mismatch:

```
1 model = CNN()
2 model.to(device)
3
4 for images, labels in train_loader:
5     images = images.to(device)
6     # labels not moved to device!
7
8     outputs = model(images)
9     loss = criterion(outputs, labels) # Error here!
```

Error type: \_\_\_\_\_

Fix: \_\_\_\_\_

#### Bug 3 - Gradient Accumulation:

```
1 for epoch in range(num_epochs):
2     for batch in train_loader:
3         outputs = model(batch)
4         loss = criterion(outputs, targets)
5         loss.backward()
6         # Missing something here...
7         optimizer.step()
```

Missing line: \_\_\_\_\_

What happens without it: \_\_\_\_\_