

# CENG403 - Spring 2025: Homework set THE-2 Study Guide

Your Name

## TASK 1: DEFORMABLE CNN MEMORIZATION GUIDE

.1

### Problem 1.1

What are the 4 corner positions for bilinear interpolation given fractional position  $q = (q_x, q_y)$ ? **Pattern to Remember:** Floor-Floor, Ceil-Floor, Floor-Ceil, Ceil-Ceil

$$p_{lt} = (\lfloor q_x \rfloor, \lfloor q_y \rfloor) \quad (\text{left top}) \quad (1)$$

$$p_{rt} = (\lceil q_x \rceil, \lfloor q_y \rfloor) \quad (\text{right top}) \quad (2)$$

$$p_{lb} = (\lfloor q_x \rfloor, \lceil q_y \rceil) \quad (\text{left bottom}) \quad (3)$$

$$p_{rb} = (\lceil q_x \rceil, \lceil q_y \rceil) \quad (\text{right bottom}) \quad (4)$$

.2

### Problem 1.2

Complete the bilinear interpolation weight formula:

$$G(p, q) = (1 - |p_x - q_x|) \cdot (1 - |p_y - q_y|)$$

.3

### Problem 1.3

Fill in the missing code for bilinear interpolation bounds checking:

```
def get_pixel_value(img, y, x):
    if 0 <= y < H and 0 <= x < W:
        return img[y, x]
    else:
        return ___ # What goes here?
```

**Answer:** 0.0 (zero padding for out-of-bounds)

.4

### Problem 1.4

What is the correct order for bilinear interpolation calculation? **Memory Pattern:** "First X, then Y"

1. Get 4 corner values:  $v_{00}, v_{01}, v_{10}, v_{11}$
2. Calculate fractional parts:  $dx = q_x - x_0, dy = q_y - y_0$
3. Interpolate along X:  $v_0 = v_{00}(1 - dx) + v_{01} \cdot dx$
4. Interpolate along X:  $v_1 = v_{10}(1 - dx) + v_{11} \cdot dx$
5. Interpolate along Y:  $out = v_0(1 - dy) + v_1 \cdot dy$

.5

## Problem 1.5

In deformable convolution, how do you extract the y and x offsets from the delta tensor? **Critical Pattern - PyTorch stores Y first, then X:**

```
delta_y = delta[n, 2 * k, h_out, w_out]    # y offset
delta_x = delta[n, 2 * k + 1, h_out, w_out] # x offset
```

.6

## Problem 1.6

What is the deformable convolution sampling position formula?

```
sample_y = h_start + kh * dilation + _____
sample_x = w_start + kw * dilation + _____
```

**Answer:** delta\_y and delta\_x

## TASK 2: CNN PYTORCH MEMORIZATION GUIDE

.1

### Problem 2.1

What are the CIFAR-100 normalization values you must memorize? **Critical Constants:**

```
mean=[0.5071, 0.4867, 0.4408] # CIFAR100 mean
std=[0.2675, 0.2565, 0.2761]  # CIFAR100 std
```

.2

### Problem 2.2

Complete the data augmentation transforms for training:

```
transform_train = transforms.Compose([
    transforms._____(32, padding=4), # What goes here?
    transforms._____,               # What goes here?
    transforms.ToTensor(),
    transforms.Normalize(mean=[...], std=[...])
])
```

**Answer:** RandomCrop and RandomHorizontalFlip

.3

### Problem 2.3

How do you split CIFAR-100 training data into 80/20 train/validation? **Pattern to Remember:**

```
train_size = int(0.8 * len(full_train_set))
val_size = len(full_train_set) - train_size
train_set, val_set = random_split(full_train_set, [train_size, val_size])
```

.4

### Problem 2.4

What is the CNN architecture pattern for the CustomCNN class? **Layer Sequence Pattern:**

1. Conv2d(3, 32) → Conv2d(32, 64) → MaxPool2d
2. Conv2d(64, 128) → MaxPool2d
3. Conv2d(128, 256) → MaxPool2d
4. Flatten → FC(256\*4\*4, 512) → FC(512, 256) → FC(256, 100)

.5

### Problem 2.5

Complete the forward pass activation pattern:

```

x = F.relu(self.conv1(x))
x = F.relu(self.conv2(x))
x = self.pool(x) # 32x32 -> 16x16
x = F.relu(self.conv3(x))
x = self.pool(x) # 16x16 -> 8x8
x = F.relu(self.conv4(x))
x = self.pool(x) # 8x8 -> 4x4
x = x.view(x.size(0), -1) # Flatten
x = F.relu(self.fc1(x))
x = self.dropout(x)
x = F.relu(self.fc2(x))
x = self.dropout(x)
x = self.fc3(x) # No activation on final layer!

```

.6

## Problem 2.6

What loss function and optimizer setup is standard for CIFAR-100?

```

loss_function = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(),
                        lr=0.01,
                        momentum=0.9,
                        weight_decay=5e-4)

```

.7

## Problem 2.7

Complete the top-1 and top-5 accuracy calculation:

```

# Top-1 accuracy
_, top1_pred = outputs.topk(1, dim=1, largest=True, sorted=True)
top1_correct = top1_pred.eq(labels.view(-1, 1)).sum().item()

# Top-5 accuracy
_, top5_pred = outputs.topk(5, dim=1, largest=True, sorted=True)
top5_correct = top5_pred.eq(labels.view(-1, 1).____(____)).sum().item()

```

Answer: `expand_as(top5_pred)`

.8

## Problem 2.8

What is the training loop structure pattern? **Memory Pattern - "Zero, Forward, Backward, Step"**:

```

optimizer.zero_grad() # Clear gradients
outputs = model(images) # Forward pass
loss = loss_function(outputs, labels) # Compute loss
loss.backward() # Backward pass
optimizer.step() # Update weights

```

.9

## Problem 2.9

How do you add BatchNorm2d to the CNN architecture? **Pattern - After each Conv2d:**

```
self.conv1 = nn.Conv2d(3, 32, kernel_size=3, padding=1)
self.bn1 = nn.BatchNorm2d(32) # Same number as conv output
self.conv2 = nn.Conv2d(32, 64, kernel_size=3, padding=1)
self.bn2 = nn.BatchNorm2d(64) # Same number as conv output
```

## TASK 3: RNN MEMORIZATION GUIDE

.1

### Problem 3.1

How do you create character vocabulary and mappings? **Standard Pattern:**

```
chars = sorted(list(set(text)))
char2idx = {ch: i for i, ch in enumerate(chars)}
idx2char = {i: ch for i, ch in enumerate(chars)}
```

.2

### Problem 3.2

How do you create input and target sequences for character prediction?

```
input_seq = text[:-1] # All except last
target_seq = text[1:] # All except first
```

.3

### Problem 3.3

Complete the one-hot encoding function:

```
def one_hot(idx, size):
    vec = torch.zeros(size)
    vec[___] = 1.0
    return vec
```

**Answer:** idx

.4

### Problem 3.4

What are the RNN weight matrix dimensions? **Dimension Memory Pattern:**

```
W_xh = torch.randn(H, V, requires_grad=True) * 0.1 # (H, V)
W_hh = torch.randn(H, H, requires_grad=True) * 0.1 # (H, H)
b_xh = torch.zeros(H, requires_grad=True)           # (H,)
b_hh = torch.zeros(H, requires_grad=True)           # (H,)
W_hy = torch.randn(V, H, requires_grad=True) * 0.1 # (V, H)
b_y = torch.zeros(V, requires_grad=True)            # (V,)
```

.5

### Problem 3.5

Complete the RNN forward pass equations:

```
# Hidden state update
h = torch.tanh(W_xh @ x_t + b_xh + W_hh @ h + b_hh)
```

```
# Output logits
s_t = ----- @ h + -----
```

**Answer:**  $W_{hy}$  and  $b_y$

.6

## Problem 3.6

How do you compute gradients explicitly with `torch.autograd`?

```
grad_W_xh = torch.autograd.grad(loss_manual, W_xh, retain_graph=True)[0]
grad_W_hh = torch.autograd.grad(loss_manual, W_hh, retain_graph=True)[0]
grad_b_xh = torch.autograd.grad(loss_manual, b_xh, retain_graph=True)[0]
grad_b_hh = torch.autograd.grad(loss_manual, b_hh, retain_graph=True)[0]
grad_W_hy = torch.autograd.grad(loss_manual, W_hy, retain_graph=True)[0]
grad_b_y = torch.autograd.grad(loss_manual, b_y, retain_graph=True)[0]
```

# CRITICAL MEMORIZATION PATTERNS

Patterns

## Problem Patterns

What are the key patterns you must memorize?

### 1. Bilinear Interpolation Pattern:

- Get 4 corners (floor/ceil combinations)
- Interpolate X first, then Y
- Use fractional parts:  $dx = q_x - x_0$ ,  $dy = q_y - y_0$

### 2. CNN Architecture Pattern:

- Conv-Conv-Pool, Conv-Pool, Conv-Pool structure
- Channel progression:  $3 \rightarrow 32 \rightarrow 64 \rightarrow 128 \rightarrow 256$
- Spatial reduction:  $32 \rightarrow 16 \rightarrow 8 \rightarrow 4$
- FC layers:  $4096 \rightarrow 512 \rightarrow 256 \rightarrow 100$

### 3. Training Loop Pattern:

- "Zero-Forward-Backward-Step"
- Always move tensors to device
- Use `torch.no_grad()` for validation

### 4. RNN Equations Pattern:

- $h_t = \tanh(W_{xh}x_t + b_{xh} + W_{hh}h_{t-1} + b_{hh})$
- $s_t = W_{hy}h_t + b_y$
- Always remember matrix dimensions

### 5. Data Preprocessing Patterns:

- CIFAR-100: `mean=[0.5071, 0.4867, 0.4408]`, `std=[0.2675, 0.2565, 0.2761]`
- 80/20 split: `train_size = int(0.8 * len(dataset))`
- Character sequences: `input = text[:-1]`, `target = text[1:]`

# COMMON MISTAKES TO AVOID

Mistakes

## Problem Mistakes

What are the most common implementation mistakes?

### 1. Deformable CNN:



- Forget PyTorch stores Y offset first, then X offset
- Wrong bilinear interpolation order (do X first, then Y)
- Forget zero padding for out-of-bounds pixels

## 2. CNN:

- Forget to move tensors to device
- Wrong flatten calculation: `x.view(x.size(0), -1)`
- Forget activation on hidden layers, add activation on output layer

## 3. RNN:

- Wrong weight matrix dimensions
- Forget `requires_grad=True` for parameters
- Use `retain_graph=True` for multiple gradient computations

# RAPID-FIRE ACTIVE RECALL QUIZ

Speed Round 1: Fill the Blanks

## Problem Speed Round 1: Fill the Blanks

Complete these critical code snippets:

**Q1:** Bilinear interpolation corners:

```
y0 = int(np._____(q_y))
x0 = int(np._____(q_x))
y1 = y0 + ____
x1 = x0 + ____
```

**Q2:** Deformable conv offset extraction:

```
delta_y = delta[n, ____ * k, h_out, w_out]
delta_x = delta[n, ____ * k + ____, h_out, w_out]
```

**Q3:** CNN flatten operation:

```
x = x.view(x.____(_), ____)
```

**Q4:** Top-5 accuracy calculation:

```
_, top5_pred = outputs.topk(____, dim=1, largest=____, sorted=____)
top5_correct = top5_pred.eq(labels.view(-1, 1)._____(____)).sum().item()
```

**Q5:** RNN hidden state update:

```
h = torch.tanh(____ @ x_t + ____ + ____ @ h + ____)
```

Speed Round 2: True/False

## Problem Speed Round 2: True/False

Mark T/F for these statements:

1. In bilinear interpolation, you interpolate Y direction first, then X direction. **[T/F]**
2. PyTorch stores Y offset before X offset in deformable convolution. **[T/F]**
3. CIFAR-100 has 100 classes, so the final FC layer outputs 100 values. **[T/F]**
4. You should apply ReLU activation to the final output layer in classification. **[T/F]**
5. In RNN,  $W_{xh}$  has dimensions (V, H). **[T/F]**
6. For validation, you need to call `optimizer.zero_grad()`. **[T/F]**
7. BatchNorm2d should be applied before the activation function. **[T/F]**
8. The input sequence for RNN is `text[1:]` and target is `text[:-1]`. **[T/F]**

### Speed Round 3: Memory Palace

## Problem Speed Round 3: Memory Palace

Associate these concepts with memorable phrases:

**Bilinear Interpolation:** "Four corners, X then Y, fractional magic"

- 4 corners: lt, rt, lb, rb
- X interpolation: top\_edge, bottom\_edge
- Y interpolation: final result

**CNN Architecture:** "3 to 32, double-double-double, then shrink to 100"

- Channels:  $3 \rightarrow 32 \rightarrow 64 \rightarrow 128 \rightarrow 256$
- Spatial:  $32 \rightarrow 16 \rightarrow 8 \rightarrow 4$
- FC:  $4096 \rightarrow 512 \rightarrow 256 \rightarrow 100$

**Training Loop:** "Zero-Forward-Backward-Step dance"

- `optimizer.zero_grad()`
- `outputs = model(images)`
- `loss.backward()`
- `optimizer.step()`

**RNN Weights:** "Input-Hidden-Hidden, Hidden-Hidden-Hidden, Hidden-Vocab-Vocab"

- $W_{xh}$ : (H, V) - maps input to hidden
- $W_{hh}$ : (H, H) - maps previous hidden to current hidden
- $W_{hy}$ : (V, H) - maps hidden to output vocabulary

## LAST-MINUTE CHECKLIST

### Pre-Exam Checklist

## Problem Pre-Exam Checklist

Before the exam, ensure you can write from memory:

**Critical Constants:**

- CIFAR-100 mean: [0.5071, 0.4867, 0.4408]
- CIFAR-100 std: [0.2675, 0.2565, 0.2761]
- Train/val split:  $0.8 * \text{len}(\text{dataset})$

**Key Formulas:**

- Bilinear weight:  $(1 - |p_x - q_x|) \cdot (1 - |p_y - q_y|)$
- RNN hidden:  $h_t = \tanh(W_{xh}x_t + b_{xh} + W_{hh}h_{t-1} + b_{hh})$
- RNN output:  $s_t = W_{hy}h_t + b_y$

#### **Critical Code Patterns:**

- Device transfer: `tensor.to(device)`
- Gradient computation: `torch.autograd.grad(loss, param, retain_graph=True)[0]`
- Top-k accuracy: `outputs.topk(k, dim=1, largest=True, sorted=True)`
- One-hot encoding: `vec[idx] = 1.0`

#### **Architecture Patterns:**

- CNN: Conv→BN→ReLU→Pool pattern
- RNN: Input→Hidden→Output with recurrence
- Training: Zero→Forward→Backward→Step

# CODING BLOCKS MEMORIZATION

Code Block 1: Bilinear Interpolation Core

## Problem Code Block 1: Bilinear Interpolation Core

Complete this bilinear interpolation function - focus on the mathematical pattern:

```
def bilinear_interpolate(a_l, q_y, q_x):
    H, W = a_l.shape

    # Step 1: Get integer positions
    y0 = int(np._____(q_y))
    x0 = int(np._____(q_x))
    y1 = y0 + ___
    x1 = x0 + ___

    # Step 2: Get values with bounds checking
    def get_pixel_value(img, y, x):
        if 0 <= y < H and 0 <= x < W:
            return img[y, x]
        else:
            return ____ # Out of bounds value

    # Step 3: Get four corner values
    v_00 = get_pixel_value(a_l, y0, x0) # ____-____
    v_01 = get_pixel_value(a_l, y0, x1) # ____-____
    v_10 = get_pixel_value(a_l, y1, x0) # ____-____
    v_11 = get_pixel_value(a_l, y1, x1) # ____-____

    # Step 4: Calculate fractional parts
    dy = q_y - ___
    dx = q_x - ___

    # Step 5: Interpolate X first, then Y
    v_0 = v_00 * (1 - ___) + v_01 * ___ # top edge
    v_1 = v_10 * (1 - ___) + v_11 * ___ # bottom edge
    out = v_0 * (1 - ___) + v_1 * ___ # final Y interpolation

    return out
```

Code Block 2: Deformable Conv Key Loop

## Problem Code Block 2: Deformable Conv Key Loop

This is the heart of deformable convolution - memorize the offset extraction pattern:

```
for n in range(N): # batch
    for c_out in range(C_out): # output channels
        for h_out in range(H_out): # height
            for w_out in range(W_out): # width
                h_start = h_out * ____
                w_start = w_out * ____
                value = 0.0

                for kh in range(K_h):
                    for kw in range(K_w):
```

```

k = kh * K_w + kw

# CRITICAL: PyTorch offset order
delta_y = delta[n, ___ * k, h_out, w_out]
delta_x = delta[n, ___ * k + ___, h_out, w_out]
m_k = mask[n, k, h_out, w_out]

# Sampling position
sample_y = h_start + kh * dilation + ____
sample_x = w_start + kw * dilation + ____

for c_in in range(C_in):
    interpolated = bilinear_interpolate(
        a_l[n, c_in, :, :], sample_y, sample_x
    )
    value += weight[c_out, c_in, kh, kw] * ___ * ____

out[n, c_out, h_out, w_out] = value

```

Code Block 3: CNN Architecture Constructor

## Problem Code Block 3: CNN Architecture Constructor

Memorize the channel progression and layer naming pattern:

```

class CustomCNN(nn.Module):
    def __init__(self, norm_layer=None):
        super(CustomCNN, self).__init__()

        # Conv layers - memorize the channel progression
        self.conv1 = nn.Conv2d(____, ___, kernel_size=3, padding=1)
        self.conv2 = nn.Conv2d(____, ___, kernel_size=3, padding=1)
        self.conv3 = nn.Conv2d(____, ___, kernel_size=3, padding=1)
        self.conv4 = nn.Conv2d(____, ___, kernel_size=3, padding=1)

        # Pooling layer
        self.pool = nn._____(2, 2)

        # FC layers - calculate the input size
        self.fc1 = nn.Linear(___ * ___ * ___, 512)
        self.fc2 = nn.Linear(____, ___)
        self.fc3 = nn.Linear(____, ___) # CIFAR-100 classes

        self.dropout = nn.Dropout(____)

```

Code Block 4: CNN Forward Pass Pattern

## Problem Code Block 4: CNN Forward Pass Pattern

Memorize the activation and pooling pattern:

```

def forward(self, x):
    # Block 1: Conv-Conv-Pool
    x = F.____(self.conv1(x))
    x = F.____(self.conv2(x))
    x = self.pool(x) # 32x32 -> ____

```

```

# Block 2: Conv-Pool
x = F.____(self.conv3(x))
x = self.pool(x) # 16x16 -> ____

# Block 3: Conv-Pool
x = F.____(self.conv4(x))
x = self.pool(x) # 8x8 -> ____

# Flatten
x = x.view(x.____(__), __)

# FC layers with dropout
x = F.____(self.fc1(x))
x = self.____(x)
x = F.____(self.fc2(x))
x = self.____(x)
x = self.fc3(x) # No activation here!

return x

```

Code Block 5: Training Loop Core

## Problem Code Block 5: Training Loop Core

The sacred training loop pattern - memorize the order:

```

def train(model, train_loader, optimizer, loss_function, device):
    model.____() # Set to training mode

    for batch in train_loader:
        images, labels = batch
        images = images.to(____)
        labels = labels.to(____)

        # The sacred four steps:
        optimizer.____() # Step 1: Clear gradients
        outputs = model(____) # Step 2: Forward pass
        loss = loss_function(outputs, labels) # Step 3: Compute loss
        loss.____() # Step 4: Backward pass
        optimizer.____() # Step 5: Update weights

        # Accuracy calculation
        _, top1_pred = outputs.topk(____, dim=1, largest=True, sorted=True)
        top1_correct = top1_pred.eq(labels.view(____, ____)).sum().item()

```

Code Block 6: BatchNorm CNN Constructor

## Problem Code Block 6: BatchNorm CNN Constructor

Pattern for adding BatchNorm after each conv layer:

```

class CustomCNNwithBN(nn.Module):
    def __init__(self, norm_layer=None):
        super(CustomCNNwithBN, self).__init__()

        self.conv1 = nn.Conv2d(3, 32, kernel_size=3, padding=1)
        self.bn1 = nn.BatchNorm2d(____) # Same as conv1 output

```

```

self.conv2 = nn.Conv2d(32, 64, kernel_size=3, padding=1)
self.bn2 = nn.BatchNorm2d(____) # Same as conv2 output

self.conv3 = nn.Conv2d(64, 128, kernel_size=3, padding=1)
self.bn3 = nn.BatchNorm2d(____) # Same as conv3 output

self.conv4 = nn.Conv2d(128, 256, kernel_size=3, padding=1)
self.bn4 = nn.BatchNorm2d(____) # Same as conv4 output

def forward(self, x):
    # Pattern: Conv -> BN -> ReLU
    x = F.relu(self.bn1(self.conv1(x)))
    x = F.relu(self.bn2(self.conv2(x)))
    x = self.pool(x)

    x = F.relu(self.bn3(self.conv3(x)))
    x = self.pool(x)

    x = F.relu(self.bn4(self.conv4(x)))
    x = self.pool(x)
    # ... rest of forward pass

```

Code Block 7: RNN Parameter Initialization

## Problem Code Block 7: RNN Parameter Initialization

Memorize the weight dimensions and initialization pattern:

```

# RNN parameters - memorize the dimensions!
W_xh = torch.randn(____, ____, requires_grad=True) * 0.1 # Input to hidden
W_hh = torch.randn(____, ____, requires_grad=True) * 0.1 # Hidden to hidden
b_xh = torch.zeros(____, requires_grad=True)             # Hidden bias 1
b_hh = torch.zeros(____, requires_grad=True)             # Hidden bias 2
W_hy = torch.randn(____, ____, requires_grad=True) * 0.1 # Hidden to output
b_y = torch.zeros(____, requires_grad=True)              # Output bias

# Remember: V = vocab size, H = hidden size

```

Code Block 8: RNN Forward Pass

## Problem Code Block 8: RNN Forward Pass

The RNN equations in code form:

```

logits_list = []
h = torch.zeros(H)

for t in range(seq_len):
    x_t = inputs[t]

    # Hidden state update equation
    h = torch.tanh(____ @ x_t + ____ + ____ @ h + ____ )

    # Output logits equation
    s_t = ____ @ h + ____

```



```

logits_list.append(s_t)

logits = torch.stack(logits_list)
log_probs = F.log_softmax(logits, dim=1)
loss_manual = F.nll_loss(log_probs, targets)

```

Code Block 9: Gradient Computation

## Problem Code Block 9: Gradient Computation

Pattern for explicit gradient computation:

```

# Compute gradients explicitly
grad_W_xh = torch.autograd.grad(loss_manual, W_xh, retain_graph=____)[0]
grad_W_hh = torch.autograd.grad(loss_manual, W_hh, retain_graph=____)[0]
grad_b_xh = torch.autograd.grad(loss_manual, b_xh, retain_graph=____)[0]
grad_b_hh = torch.autograd.grad(loss_manual, b_hh, retain_graph=____)[0]
grad_W_hy = torch.autograd.grad(loss_manual, W_hy, retain_graph=____)[0]
grad_b_y = torch.autograd.grad(loss_manual, b_y, retain_graph=____)[0]

# Why retain_graph=True? Because we compute multiple gradients from same loss

```

Code Block 10: Character Processing

## Problem Code Block 10: Character Processing

Standard pattern for character-level RNN preprocessing:

```

text = "Deep Learning"

# Step 1: Create vocabulary
chars = sorted(list(set(____)))
char2idx = {ch: i for i, ch in enumerate(____)}
idx2char = {i: ch for i, ch in enumerate(____)}

# Step 2: Create sequences
input_seq = text[____] # All except last
target_seq = text[____] # All except first

# Step 3: Convert to tensors
inputs = [one_hot(char2idx[ch], V) for ch in ____]
targets = torch.tensor([char2idx[ch] for ch in ____], dtype=torch.long)

def one_hot(idx, size):
    vec = torch.zeros(size)
    vec[____] = 1.0
    return vec

```

# MEMORIZATION MNEMONICS

Memory Aids

## Problem Memory Aids

Use these phrases to remember key patterns:

**”Floor-Ceil-Four-Corners”**: Bilinear interpolation corners

- lt: floor-floor, rt: ceil-floor, lb: floor-ceil, rb: ceil-ceil

**”Y-before-X-in-PyTorch”**: Deformable conv offset ordering

- $\text{delta\_y} = \text{delta}[n, 2*k, \text{h\_out}, \text{w\_out}]$
- $\text{delta\_x} = \text{delta}[n, 2*k+1, \text{h\_out}, \text{w\_out}]$

**”3-32-64-128-256”**: CNN channel progression

- Each layer doubles the channels (except first)

**”32-16-8-4”**: Spatial dimension reduction

- Each `MaxPool2d(2,2)` halves the spatial dimensions

**”Zero-Forward-Backward-Step”**: Training loop mantra

- Never forget the order!

**”Input-Hidden-Hidden”**: RNN weight dimensions

- $W_{xh}$ : (H,V),  $W_{hh}$ : (H,H),  $W_{hy}$ : (V,H)

**”Tanh-Hidden-Linear-Output”**: RNN computation flow

- Hidden uses tanh, output is linear

## ANSWERS TO CODING BLOCKS

**Code Block 1**: floor, floor, 1, 1, 0.0, top-left, top-right, bottom-left, bottom-right, y0, x0, dx, dx, dx, dx, dy, dy

**Code Block 2**: stride, stride, 2, 2, 1, delta\_y, delta\_x, m\_k, interpolated

**Code Block 3**: 3, 32, 32, 64, 64, 128, 128, 256, `MaxPool2d`, 256, 4, 4, 512, 256, 256, 100, 0.5

**Code Block 4**: relu, relu, 16x16, relu, 8x8, relu, 4x4, `size(0)`, -1, relu, dropout, relu, dropout

**Code Block 5**: train, device, device, zero\_grad, images, backward, step, 1, -1, 1

**Code Block 6:** 32, 64, 128, 256

**Code Block 7:** H, V, H, H, H, H, V, H, V

**Code Block 8:** W\_xh, b\_xh, W\_hh, b\_hh, W\_hy, b\_y

**Code Block 9:** True, True, True, True, True, True

# ADVANCED CODING SCENARIOS

## Problem 1

Scenario 1: Debugging Deformable Conv If your deformable convolution gives wrong results, what are the most likely bugs?

```
# Common Bug 1: Wrong offset extraction
delta_y = delta[n, k, h_out, w_out]      # WRONG - missing factor of 2
delta_x = delta[n, k + 1, h_out, w_out]   # WRONG - should be 2*k+1

# Correct version:
delta_y = delta[n, ____ * k, h_out, w_out]
delta_x = delta[n, ____ * k + ____, h_out, w_out]

# Common Bug 2: Wrong bilinear interpolation order
# WRONG: Interpolate Y first
v_y = v_00 * (1 - dy) + v_10 * dy
v_final = v_y * (1 - dx) + v_01 * dx

# Correct: Interpolate X first, then Y
v_0 = v_00 * (1 - ____ ) + v_01 * ____   # top edge
v_1 = v_10 * (1 - ____ ) + v_11 * ____   # bottom edge
out = v_0 * (1 - ____ ) + v_1 * ____      # Y interpolation
```

Scenario 2: CNN Architecture Variations

## Problem Scenario 2: CNN Architecture Variations

If asked to modify the CNN, remember these patterns:

```
# Adding more conv layers - maintain the pattern
class CustomCNNDeep(nn.Module):
    def __init__(self):
        super().__init__()
        # Pattern: start with 3 channels, double each time
        self.conv1 = nn.Conv2d(3, 32, 3, padding=1)
        self.conv2 = nn.Conv2d(32, 32, 3, padding=1) # Same channels
        self.conv3 = nn.Conv2d(32, 64, 3, padding=1) # Double
        self.conv4 = nn.Conv2d(64, 64, 3, padding=1) # Same
        self.conv5 = nn.Conv2d(64, 128, 3, padding=1) # Double

    def forward(self, x):
        # Pattern: Conv-Conv-Pool, Conv-Conv-Pool
        x = F.relu(self.conv1(x))
        x = F.relu(self.conv2(x))
        x = self.pool(x) # After every 2 conv layers

        x = F.relu(self.conv3(x))
        x = F.relu(self.conv4(x))
        x = self.pool(x)

        # Calculate new flatten size: 128 * 8 * 8 = ____
```

Scenario 3: Validation vs Training Mode

## Problem Scenario 3: Validation vs Training Mode

Critical differences between training and validation:

```
# Training mode
def train_epoch():
    model.____() # Enable dropout and batch norm training mode

    for batch in train_loader:
        optimizer.____() # Clear gradients
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.____() # Compute gradients
        optimizer.____() # Update weights

# Validation mode
def validate():
    model.____() # Disable dropout, batch norm in eval mode

    with torch.____(): # Disable gradient computation
        for batch in val_loader:
            # NO optimizer.zero_grad() here!
            # NO loss.backward() here!
            # NO optimizer.step() here!
            outputs = model(images)
            loss = criterion(outputs, labels)
```

Scenario 4: RNN with Different Sequence Lengths

## Problem Scenario 4: RNN with Different Sequence Lengths

If given a different text, adapt the RNN code:

```
# Original: "Deep Learning"
text = "Deep Learning"
input_seq = text[:-1] # "Deep Learnin"
target_seq = text[1:] # "eep Learning"

# New text: "Hello World"
text = "Hello World"
input_seq = text[____] # "Hello Worl"
target_seq = text[____] # "ello World"

# Vocabulary size changes!
chars = sorted(list(set(text)))
V = len(chars) # This will be different!

# All weight matrices need to be reinitialized with new V
W_xh = torch.randn(H, ____, requires_grad=True) * 0.1
W_hy = torch.randn(____, H, requires_grad=True) * 0.1
b_y = torch.zeros(____, requires_grad=True)
```

Scenario 5: Hyperparameter Grid Search Pattern

## Problem Scenario 5: Hyperparameter Grid Search Pattern

Standard grid search implementation:

```
learning_rates = [0.0001, 0.001]
```

```

optimizers = [torch.optim.Adam, torch.optim.SGD]
model_classes = [CustomCNN, CustomCNNwithBN]

best_accuracy = 0
best_params = None

for model_class in model_classes:
    for optimizer_class in optimizers:
        for lr in learning_rates:
            # CRITICAL: Reinitialize model each time
            model = model_class().to(device)

            # Initialize optimizer based on type
            if optimizer_class == torch.optim.SGD:
                optimizer = optimizer_class(
                    model.parameters(),
                    lr=lr,
                    momentum=____,
                    weight_decay=____
                )
            else: # Adam
                optimizer = optimizer_class(
                    model.parameters(),
                    lr=lr,
                    weight_decay=____
                )

            # Train and validate...
            val_acc = train_and_validate(model, optimizer)

            if val_acc > best_accuracy:
                best_accuracy = val_acc
                best_params = (model_class.__name__, optimizer_class.__name__, lr)

```

Scenario 6: Top-K Accuracy Calculation Variations

## Problem Scenario 6: Top-K Accuracy Calculation Variations

Different ways to calculate accuracy:

```

# Top-1 accuracy (most common)
_, predicted = torch.max(outputs, 1)
correct = (predicted == labels).sum().item()
accuracy = correct / labels.size(0) * 100

# Top-K accuracy using topk
_, top_k_pred = outputs.topk(____, dim=1, largest=True, sorted=True)
top_k_correct = top_k_pred.eq(labels.view(-1, 1).expand_as(____)).sum().item()

# Alternative top-K calculation
values, indices = torch.topk(outputs, k=5, dim=1)
correct_mask = indices == labels.unsqueeze(1)
top_k_accuracy = correct_mask.sum().float() / labels.size(0) * 100

```

# EXAM SIMULATION QUESTIONS

Quick Code Writing

## Problem Quick Code Writing

Write these functions from memory in 2 minutes each:

1. Write the bilinear interpolation weight calculation:

```
def calculate_bilinear_weight(p_x, p_y, q_x, q_y):  
    # Your code here - calculate G(p,q)  
    return ____
```

2. Write the RNN hidden state update:

```
def rnn_step(x_t, h_prev, W_xh, W_hh, b_xh, b_hh):  
    # Your code here - compute new hidden state  
    return ____
```

3. Write the CNN forward pass for one block:

```
def cnn_block_forward(x, conv1, conv2, pool):  
    # Your code here - conv-conv-pool pattern  
    return ____
```

4. Write the training step:

```
def training_step(model, optimizer, criterion, images, labels):  
    # Your code here - complete training step  
    return loss
```

5. Write character to one-hot conversion:

```
def char_to_onehot(char, char2idx, vocab_size):  
    # Your code here - convert character to one-hot vector  
    return ____
```

## FINAL MEMORY CHECK

### Problem 1

Last Minute Review Before the exam, quickly verify you remember:

Constants (write from memory):

- CIFAR-100 mean: [\_\_\_\_, \_\_\_\_, \_\_\_\_]
- CIFAR-100 std: [\_\_\_\_, \_\_\_\_, \_\_\_\_]
- Train/val split ratio: \_\_\_\_

- SGD momentum: \_\_\_\_

Dimensions (write from memory):

- RNN W\_hy: (----, ----)

**Key Equations (write from memory):**

- RNN output:  $s_t = \text{-----}$



# CENG403 - Spring 2025: Homework set THE-2 Study Guide

Your Name

## TASK 1: DEFORMABLE CNN MEMORIZATION GUIDE

.1

### Problem 1.1

What are the 4 corner positions for bilinear interpolation given fractional position  $q = (q_x, q_y)$ ? **Pattern to Remember:** Floor-Floor, Ceil-Floor, Floor-Ceil, Ceil-Ceil

$$p_{lt} = (\lfloor q_x \rfloor, \lfloor q_y \rfloor) \quad (\text{left top}) \quad (1)$$

$$p_{rt} = (\lceil q_x \rceil, \lfloor q_y \rfloor) \quad (\text{right top}) \quad (2)$$

$$p_{lb} = (\lfloor q_x \rfloor, \lceil q_y \rceil) \quad (\text{left bottom}) \quad (3)$$

$$p_{rb} = (\lceil q_x \rceil, \lceil q_y \rceil) \quad (\text{right bottom}) \quad (4)$$

.2

### Problem 1.2

Complete the bilinear interpolation weight formula:

$$G(p, q) = (1 - |p_x - q_x|) \cdot (1 - |p_y - q_y|)$$

.3

### Problem 1.3

Fill in the missing code for bilinear interpolation bounds checking:

```
def get_pixel_value(img, y, x):
    if 0 <= y < H and 0 <= x < W:
        return img[y, x]
    else:
        return ___ # What goes here?
```

**Answer:** 0.0 (zero padding for out-of-bounds)

.4

### Problem 1.4

What is the correct order for bilinear interpolation calculation? **Memory Pattern:** "First X, then Y"

1. Get 4 corner values:  $v_{00}, v_{01}, v_{10}, v_{11}$
2. Calculate fractional parts:  $dx = q_x - x_0, dy = q_y - y_0$
3. Interpolate along X:  $v_0 = v_{00}(1 - dx) + v_{01} \cdot dx$
4. Interpolate along X:  $v_1 = v_{10}(1 - dx) + v_{11} \cdot dx$
5. Interpolate along Y:  $out = v_0(1 - dy) + v_1 \cdot dy$

.5

## Problem 1.5

In deformable convolution, how do you extract the y and x offsets from the delta tensor? **Critical Pattern - PyTorch stores Y first, then X:**

```
delta_y = delta[n, 2 * k, h_out, w_out]    # y offset
delta_x = delta[n, 2 * k + 1, h_out, w_out] # x offset
```

.6

## Problem 1.6

What is the deformable convolution sampling position formula?

```
sample_y = h_start + kh * dilation + _____
sample_x = w_start + kw * dilation + _____
```

**Answer:** delta\_y and delta\_x

## TASK 2: CNN PYTORCH MEMORIZATION GUIDE

.1

### Problem 2.1

What are the CIFAR-100 normalization values you must memorize? **Critical Constants:**

```
mean=[0.5071, 0.4867, 0.4408] # CIFAR100 mean
std=[0.2675, 0.2565, 0.2761]  # CIFAR100 std
```

.2

### Problem 2.2

Complete the data augmentation transforms for training:

```
transform_train = transforms.Compose([
    transforms._____(32, padding=4), # What goes here?
    transforms._____,               # What goes here?
    transforms.ToTensor(),
    transforms.Normalize(mean=[...], std=[...])
])
```

**Answer:** RandomCrop and RandomHorizontalFlip

.3

### Problem 2.3

How do you split CIFAR-100 training data into 80/20 train/validation? **Pattern to Remember:**

```
train_size = int(0.8 * len(full_train_set))
val_size = len(full_train_set) - train_size
train_set, val_set = random_split(full_train_set, [train_size, val_size])
```

.4

### Problem 2.4

What is the CNN architecture pattern for the CustomCNN class? **Layer Sequence Pattern:**

1. Conv2d(3, 32) → Conv2d(32, 64) → MaxPool2d
2. Conv2d(64, 128) → MaxPool2d
3. Conv2d(128, 256) → MaxPool2d
4. Flatten → FC(256\*4\*4, 512) → FC(512, 256) → FC(256, 100)

.5

### Problem 2.5

Complete the forward pass activation pattern:

```

x = F.relu(self.conv1(x))
x = F.relu(self.conv2(x))
x = self.pool(x) # 32x32 -> 16x16
x = F.relu(self.conv3(x))
x = self.pool(x) # 16x16 -> 8x8
x = F.relu(self.conv4(x))
x = self.pool(x) # 8x8 -> 4x4
x = x.view(x.size(0), -1) # Flatten
x = F.relu(self.fc1(x))
x = self.dropout(x)
x = F.relu(self.fc2(x))
x = self.dropout(x)
x = self.fc3(x) # No activation on final layer!

```

.6

## Problem 2.6

What loss function and optimizer setup is standard for CIFAR-100?

```

loss_function = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(),
                        lr=0.01,
                        momentum=0.9,
                        weight_decay=5e-4)

```

.7

## Problem 2.7

Complete the top-1 and top-5 accuracy calculation:

```

# Top-1 accuracy
_, top1_pred = outputs.topk(1, dim=1, largest=True, sorted=True)
top1_correct = top1_pred.eq(labels.view(-1, 1)).sum().item()

# Top-5 accuracy
_, top5_pred = outputs.topk(5, dim=1, largest=True, sorted=True)
top5_correct = top5_pred.eq(labels.view(-1, 1).____(____)).sum().item()

```

Answer: `expand_as(top5_pred)`

.8

## Problem 2.8

What is the training loop structure pattern? **Memory Pattern - "Zero, Forward, Backward, Step"**:

```

optimizer.zero_grad() # Clear gradients
outputs = model(images) # Forward pass
loss = loss_function(outputs, labels) # Compute loss
loss.backward() # Backward pass
optimizer.step() # Update weights

```

.9

## Problem 2.9

How do you add BatchNorm2d to the CNN architecture? **Pattern - After each Conv2d:**

```
self.conv1 = nn.Conv2d(3, 32, kernel_size=3, padding=1)
self.bn1 = nn.BatchNorm2d(32) # Same number as conv output
self.conv2 = nn.Conv2d(32, 64, kernel_size=3, padding=1)
self.bn2 = nn.BatchNorm2d(64) # Same number as conv output
```

## TASK 3: RNN MEMORIZATION GUIDE

.1

### Problem 3.1

How do you create character vocabulary and mappings? **Standard Pattern:**

```
chars = sorted(list(set(text)))
char2idx = {ch: i for i, ch in enumerate(chars)}
idx2char = {i: ch for i, ch in enumerate(chars)}
```

.2

### Problem 3.2

How do you create input and target sequences for character prediction?

```
input_seq = text[:-1] # All except last
target_seq = text[1:] # All except first
```

.3

### Problem 3.3

Complete the one-hot encoding function:

```
def one_hot(idx, size):
    vec = torch.zeros(size)
    vec[___] = 1.0
    return vec
```

**Answer:** idx

.4

### Problem 3.4

What are the RNN weight matrix dimensions? **Dimension Memory Pattern:**

```
W_xh = torch.randn(H, V, requires_grad=True) * 0.1 # (H, V)
W_hh = torch.randn(H, H, requires_grad=True) * 0.1 # (H, H)
b_xh = torch.zeros(H, requires_grad=True)          # (H,)
b_hh = torch.zeros(H, requires_grad=True)          # (H,)
W_hy = torch.randn(V, H, requires_grad=True) * 0.1 # (V, H)
b_y = torch.zeros(V, requires_grad=True)           # (V,)
```

.5

### Problem 3.5

Complete the RNN forward pass equations:

```
# Hidden state update
h = torch.tanh(W_xh @ x_t + b_xh + W_hh @ h + b_hh)
```

```
# Output logits
s_t = ----- @ h + -----
```

**Answer:**  $W_{hy}$  and  $b_y$

.6

## Problem 3.6

How do you compute gradients explicitly with `torch.autograd`?

```
grad_W_xh = torch.autograd.grad(loss_manual, W_xh, retain_graph=True)[0]
grad_W_hh = torch.autograd.grad(loss_manual, W_hh, retain_graph=True)[0]
grad_b_xh = torch.autograd.grad(loss_manual, b_xh, retain_graph=True)[0]
grad_b_hh = torch.autograd.grad(loss_manual, b_hh, retain_graph=True)[0]
grad_W_hy = torch.autograd.grad(loss_manual, W_hy, retain_graph=True)[0]
grad_b_y = torch.autograd.grad(loss_manual, b_y, retain_graph=True)[0]
```

# CRITICAL MEMORIZATION PATTERNS

Patterns

## Problem Patterns

What are the key patterns you must memorize?

### 1. Bilinear Interpolation Pattern:

- Get 4 corners (floor/ceil combinations)
- Interpolate X first, then Y
- Use fractional parts:  $dx = q_x - x_0$ ,  $dy = q_y - y_0$

### 2. CNN Architecture Pattern:

- Conv-Conv-Pool, Conv-Pool, Conv-Pool structure
- Channel progression:  $3 \rightarrow 32 \rightarrow 64 \rightarrow 128 \rightarrow 256$
- Spatial reduction:  $32 \rightarrow 16 \rightarrow 8 \rightarrow 4$
- FC layers:  $4096 \rightarrow 512 \rightarrow 256 \rightarrow 100$

### 3. Training Loop Pattern:

- "Zero-Forward-Backward-Step"
- Always move tensors to device
- Use `torch.no_grad()` for validation

### 4. RNN Equations Pattern:

- $h_t = \tanh(W_{xh}x_t + b_{xh} + W_{hh}h_{t-1} + b_{hh})$
- $s_t = W_{hy}h_t + b_y$
- Always remember matrix dimensions

### 5. Data Preprocessing Patterns:

- CIFAR-100: `mean=[0.5071, 0.4867, 0.4408]`, `std=[0.2675, 0.2565, 0.2761]`
- 80/20 split: `train_size = int(0.8 * len(dataset))`
- Character sequences: `input = text[:-1]`, `target = text[1:]`

# COMMON MISTAKES TO AVOID

Mistakes

## Problem Mistakes

What are the most common implementation mistakes?

### 1. Deformable CNN:



- Forget PyTorch stores Y offset first, then X offset
- Wrong bilinear interpolation order (do X first, then Y)
- Forget zero padding for out-of-bounds pixels

## 2. CNN:

- Forget to move tensors to device
- Wrong flatten calculation: `x.view(x.size(0), -1)`
- Forget activation on hidden layers, add activation on output layer

## 3. RNN:

- Wrong weight matrix dimensions
- Forget `requires_grad=True` for parameters
- Use `retain_graph=True` for multiple gradient computations

# RAPID-FIRE ACTIVE RECALL QUIZ

Speed Round 1: Fill the Blanks

## Problem Speed Round 1: Fill the Blanks

Complete these critical code snippets:

**Q1:** Bilinear interpolation corners:

```
y0 = int(np._____(q_y))
x0 = int(np._____(q_x))
y1 = y0 + ____
x1 = x0 + ____
```

**Q2:** Deformable conv offset extraction:

```
delta_y = delta[n, ____ * k, h_out, w_out]
delta_x = delta[n, ____ * k + ____, h_out, w_out]
```

**Q3:** CNN flatten operation:

```
x = x.view(x.____(__), ____)
```

**Q4:** Top-5 accuracy calculation:

```
_, top5_pred = outputs.topk(____, dim=1, largest=____, sorted=____)
top5_correct = top5_pred.eq(labels.view(-1, 1)._____(____)).sum().item()
```

**Q5:** RNN hidden state update:

```
h = torch.tanh(____ @ x_t + _____ + _____ @ h + _____)
```

Speed Round 2: True/False

## Problem Speed Round 2: True/False

Mark T/F for these statements:

1. In bilinear interpolation, you interpolate Y direction first, then X direction. **[T/F]**
2. PyTorch stores Y offset before X offset in deformable convolution. **[T/F]**
3. CIFAR-100 has 100 classes, so the final FC layer outputs 100 values. **[T/F]**
4. You should apply ReLU activation to the final output layer in classification. **[T/F]**
5. In RNN,  $W_{xh}$  has dimensions (V, H). **[T/F]**
6. For validation, you need to call `optimizer.zero_grad()`. **[T/F]**
7. BatchNorm2d should be applied before the activation function. **[T/F]**
8. The input sequence for RNN is `text[1:]` and target is `text[:-1]`. **[T/F]**

### Speed Round 3: Memory Palace

## Problem Speed Round 3: Memory Palace

Associate these concepts with memorable phrases:

**Bilinear Interpolation:** "Four corners, X then Y, fractional magic"

- 4 corners: lt, rt, lb, rb
- X interpolation: top\_edge, bottom\_edge
- Y interpolation: final result

**CNN Architecture:** "3 to 32, double-double-double, then shrink to 100"

- Channels:  $3 \rightarrow 32 \rightarrow 64 \rightarrow 128 \rightarrow 256$
- Spatial:  $32 \rightarrow 16 \rightarrow 8 \rightarrow 4$
- FC:  $4096 \rightarrow 512 \rightarrow 256 \rightarrow 100$

**Training Loop:** "Zero-Forward-Backward-Step dance"

- `optimizer.zero_grad()`
- `outputs = model(images)`
- `loss.backward()`
- `optimizer.step()`

**RNN Weights:** "Input-Hidden-Hidden, Hidden-Hidden-Hidden, Hidden-Vocab-Vocab"

- $W_{xh}$ : (H, V) - maps input to hidden
- $W_{hh}$ : (H, H) - maps previous hidden to current hidden
- $W_{hy}$ : (V, H) - maps hidden to output vocabulary

## LAST-MINUTE CHECKLIST

### Pre-Exam Checklist

## Problem Pre-Exam Checklist

Before the exam, ensure you can write from memory:

**Critical Constants:**

- CIFAR-100 mean: [0.5071, 0.4867, 0.4408]
- CIFAR-100 std: [0.2675, 0.2565, 0.2761]
- Train/val split:  $0.8 * \text{len}(\text{dataset})$

**Key Formulas:**

- Bilinear weight:  $(1 - |p_x - q_x|) \cdot (1 - |p_y - q_y|)$
- RNN hidden:  $h_t = \tanh(W_{xh}x_t + b_{xh} + W_{hh}h_{t-1} + b_{hh})$
- RNN output:  $s_t = W_{hy}h_t + b_y$

#### Critical Code Patterns:

- Device transfer: `tensor.to(device)`
- Gradient computation: `torch.autograd.grad(loss, param, retain_graph=True)[0]`
- Top-k accuracy: `outputs.topk(k, dim=1, largest=True, sorted=True)`
- One-hot encoding: `vec[idx] = 1.0`

#### Architecture Patterns:

- CNN: Conv→BN→ReLU→Pool pattern
- RNN: Input→Hidden→Output with recurrence
- Training: Zero→Forward→Backward→Step

# CODING BLOCKS MEMORIZATION

Code Block 1: Bilinear Interpolation Core

## Problem Code Block 1: Bilinear Interpolation Core

Complete this bilinear interpolation function - focus on the mathematical pattern:

```
def bilinear_interpolate(a_l, q_y, q_x):
    H, W = a_l.shape

    # Step 1: Get integer positions
    y0 = int(np._____(q_y))
    x0 = int(np._____(q_x))
    y1 = y0 + ___
    x1 = x0 + ___

    # Step 2: Get values with bounds checking
    def get_pixel_value(img, y, x):
        if 0 <= y < H and 0 <= x < W:
            return img[y, x]
        else:
            return ____ # Out of bounds value

    # Step 3: Get four corner values
    v_00 = get_pixel_value(a_l, y0, x0) # ____-____
    v_01 = get_pixel_value(a_l, y0, x1) # ____-____
    v_10 = get_pixel_value(a_l, y1, x0) # ____-____
    v_11 = get_pixel_value(a_l, y1, x1) # ____-____

    # Step 4: Calculate fractional parts
    dy = q_y - ___
    dx = q_x - ___

    # Step 5: Interpolate X first, then Y
    v_0 = v_00 * (1 - ___) + v_01 * ___ # top edge
    v_1 = v_10 * (1 - ___) + v_11 * ___ # bottom edge
    out = v_0 * (1 - ___) + v_1 * ___ # final Y interpolation

    return out
```

Code Block 2: Deformable Conv Key Loop

## Problem Code Block 2: Deformable Conv Key Loop

This is the heart of deformable convolution - memorize the offset extraction pattern:

```
for n in range(N): # batch
    for c_out in range(C_out): # output channels
        for h_out in range(H_out): # height
            for w_out in range(W_out): # width
                h_start = h_out * ____
                w_start = w_out * ____
                value = 0.0

                for kh in range(K_h):
                    for kw in range(K_w):
```

```

k = kh * K_w + kw

# CRITICAL: PyTorch offset order
delta_y = delta[n, ___ * k, h_out, w_out]
delta_x = delta[n, ___ * k + ___, h_out, w_out]
m_k = mask[n, k, h_out, w_out]

# Sampling position
sample_y = h_start + kh * dilation + ____
sample_x = w_start + kw * dilation + ____

for c_in in range(C_in):
    interpolated = bilinear_interpolate(
        a_l[n, c_in, :, :], sample_y, sample_x
    )
    value += weight[c_out, c_in, kh, kw] * ___ * ____

out[n, c_out, h_out, w_out] = value

```

Code Block 3: CNN Architecture Constructor

## Problem Code Block 3: CNN Architecture Constructor

Memorize the channel progression and layer naming pattern:

```

class CustomCNN(nn.Module):
    def __init__(self, norm_layer=None):
        super(CustomCNN, self).__init__()

        # Conv layers - memorize the channel progression
        self.conv1 = nn.Conv2d(____, ___, kernel_size=3, padding=1)
        self.conv2 = nn.Conv2d(____, ___, kernel_size=3, padding=1)
        self.conv3 = nn.Conv2d(____, ___, kernel_size=3, padding=1)
        self.conv4 = nn.Conv2d(____, ___, kernel_size=3, padding=1)

        # Pooling layer
        self.pool = nn._____(2, 2)

        # FC layers - calculate the input size
        self.fc1 = nn.Linear(___ * ___ * ___, 512)
        self.fc2 = nn.Linear(____, ___)
        self.fc3 = nn.Linear(____, ___) # CIFAR-100 classes

        self.dropout = nn.Dropout(____)

```

Code Block 4: CNN Forward Pass Pattern

## Problem Code Block 4: CNN Forward Pass Pattern

Memorize the activation and pooling pattern:

```

def forward(self, x):
    # Block 1: Conv-Conv-Pool
    x = F.____(self.conv1(x))
    x = F.____(self.conv2(x))
    x = self.pool(x) # 32x32 -> ____

```

```

# Block 2: Conv-Pool
x = F.____(self.conv3(x))
x = self.pool(x) # 16x16 -> ____

# Block 3: Conv-Pool
x = F.____(self.conv4(x))
x = self.pool(x) # 8x8 -> ____

# Flatten
x = x.view(x.____(__), __)

# FC layers with dropout
x = F.____(self.fc1(x))
x = self.____(x)
x = F.____(self.fc2(x))
x = self.____(x)
x = self.fc3(x) # No activation here!

return x

```

Code Block 5: Training Loop Core

## Problem Code Block 5: Training Loop Core

The sacred training loop pattern - memorize the order:

```

def train(model, train_loader, optimizer, loss_function, device):
    model.____() # Set to training mode

    for batch in train_loader:
        images, labels = batch
        images = images.to(____)
        labels = labels.to(____)

        # The sacred four steps:
        optimizer.____() # Step 1: Clear gradients
        outputs = model(____) # Step 2: Forward pass
        loss = loss_function(outputs, labels) # Step 3: Compute loss
        loss.____() # Step 4: Backward pass
        optimizer.____() # Step 5: Update weights

        # Accuracy calculation
        _, top1_pred = outputs.topk(____, dim=1, largest=True, sorted=True)
        top1_correct = top1_pred.eq(labels.view(____, ____)).sum().item()

```

Code Block 6: BatchNorm CNN Constructor

## Problem Code Block 6: BatchNorm CNN Constructor

Pattern for adding BatchNorm after each conv layer:

```

class CustomCNNwithBN(nn.Module):
    def __init__(self, norm_layer=None):
        super(CustomCNNwithBN, self).__init__()

        self.conv1 = nn.Conv2d(3, 32, kernel_size=3, padding=1)
        self.bn1 = nn.BatchNorm2d(____) # Same as conv1 output

```

```

self.conv2 = nn.Conv2d(32, 64, kernel_size=3, padding=1)
self.bn2 = nn.BatchNorm2d(____) # Same as conv2 output

self.conv3 = nn.Conv2d(64, 128, kernel_size=3, padding=1)
self.bn3 = nn.BatchNorm2d(____) # Same as conv3 output

self.conv4 = nn.Conv2d(128, 256, kernel_size=3, padding=1)
self.bn4 = nn.BatchNorm2d(____) # Same as conv4 output

def forward(self, x):
    # Pattern: Conv -> BN -> ReLU
    x = F.relu(self.bn1(self.conv1(x)))
    x = F.relu(self.bn2(self.conv2(x)))
    x = self.pool(x)

    x = F.relu(self.bn3(self.conv3(x)))
    x = self.pool(x)

    x = F.relu(self.bn4(self.conv4(x)))
    x = self.pool(x)
    # ... rest of forward pass

```

Code Block 7: RNN Parameter Initialization

## Problem Code Block 7: RNN Parameter Initialization

Memorize the weight dimensions and initialization pattern:

```

# RNN parameters - memorize the dimensions!
W_xh = torch.randn(____, ____, requires_grad=True) * 0.1 # Input to hidden
W_hh = torch.randn(____, ____, requires_grad=True) * 0.1 # Hidden to hidden
b_xh = torch.zeros(____, requires_grad=True) # Hidden bias 1
b_hh = torch.zeros(____, requires_grad=True) # Hidden bias 2
W_hy = torch.randn(____, ____, requires_grad=True) * 0.1 # Hidden to output
b_y = torch.zeros(____, requires_grad=True) # Output bias

# Remember: V = vocab size, H = hidden size

```

Code Block 8: RNN Forward Pass

## Problem Code Block 8: RNN Forward Pass

The RNN equations in code form:

```

logits_list = []
h = torch.zeros(H)

for t in range(seq_len):
    x_t = inputs[t]

    # Hidden state update equation
    h = torch.tanh(____ @ x_t + ____ + ____ @ h + ____ )

    # Output logits equation
    s_t = ____ @ h + ____

```



```

logits_list.append(s_t)

logits = torch.stack(logits_list)
log_probs = F.log_softmax(logits, dim=1)
loss_manual = F.nll_loss(log_probs, targets)

```

Code Block 9: Gradient Computation

## Problem Code Block 9: Gradient Computation

Pattern for explicit gradient computation:

```

# Compute gradients explicitly
grad_W_xh = torch.autograd.grad(loss_manual, W_xh, retain_graph=____)[0]
grad_W_hh = torch.autograd.grad(loss_manual, W_hh, retain_graph=____)[0]
grad_b_xh = torch.autograd.grad(loss_manual, b_xh, retain_graph=____)[0]
grad_b_hh = torch.autograd.grad(loss_manual, b_hh, retain_graph=____)[0]
grad_W_hy = torch.autograd.grad(loss_manual, W_hy, retain_graph=____)[0]
grad_b_y = torch.autograd.grad(loss_manual, b_y, retain_graph=____)[0]

# Why retain_graph=True? Because we compute multiple gradients from same loss

```

Code Block 10: Character Processing

## Problem Code Block 10: Character Processing

Standard pattern for character-level RNN preprocessing:

```

text = "Deep Learning"

# Step 1: Create vocabulary
chars = sorted(list(set(____)))
char2idx = {ch: i for i, ch in enumerate(____)}
idx2char = {i: ch for i, ch in enumerate(____)}

# Step 2: Create sequences
input_seq = text[____] # All except last
target_seq = text[____] # All except first

# Step 3: Convert to tensors
inputs = [one_hot(char2idx[ch], V) for ch in ____]
targets = torch.tensor([char2idx[ch] for ch in ____], dtype=torch.long)

def one_hot(idx, size):
    vec = torch.zeros(size)
    vec[____] = 1.0
    return vec

```

# MEMORIZATION MNEMONICS

Memory Aids

## Problem Memory Aids

Use these phrases to remember key patterns:

**"Floor-Ceil-Four-Corners":** Bilinear interpolation corners

- lt: floor-floor, rt: ceil-floor, lb: floor-ceil, rb: ceil-ceil

**"Y-before-X-in-PyTorch":** Deformable conv offset ordering

- $\text{delta\_y} = \text{delta}[n, 2*k, \text{h\_out}, \text{w\_out}]$
- $\text{delta\_x} = \text{delta}[n, 2*k+1, \text{h\_out}, \text{w\_out}]$

**"3-32-64-128-256":** CNN channel progression

- Each layer doubles the channels (except first)

**"32-16-8-4":** Spatial dimension reduction

- Each `MaxPool2d(2,2)` halves the spatial dimensions

**"Zero-Forward-Backward-Step":** Training loop mantra

- Never forget the order!

**"Input-Hidden-Hidden":** RNN weight dimensions

- $W_{xh}$ : (H,V),  $W_{hh}$ : (H,H),  $W_{hy}$ : (V,H)

**"Tanh-Hidden-Linear-Output":** RNN computation flow

- Hidden uses tanh, output is linear

## ANSWERS TO CODING BLOCKS

**Code Block 1:** floor, floor, 1, 1, 0.0, top-left, top-right, bottom-left, bottom-right, y0, x0, dx, dx, dx, dx, dy, dy

**Code Block 2:** stride, stride, 2, 2, 1, delta\_y, delta\_x, m\_k, interpolated

**Code Block 3:** 3, 32, 32, 64, 64, 128, 128, 256, `MaxPool2d`, 256, 4, 4, 512, 256, 256, 100, 0.5

**Code Block 4:** relu, relu, 16x16, relu, 8x8, relu, 4x4, `size(0)`, -1, relu, dropout, relu, dropout

**Code Block 5:** train, device, device, zero\_grad, images, backward, step, 1, -1, 1

**Code Block 6:** 32, 64, 128, 256

**Code Block 7:** H, V, H, H, H, H, V, H, V

**Code Block 8:** W\_xh, b\_xh, W\_hh, b\_hh, W\_hy, b\_y

**Code Block 9:** True, True, True, True, True, True

# ADVANCED CODING SCENARIOS

## Scenario 1: Debugging Deformable Conv

### Problem Scenario 1: Debugging Deformable Conv

If your deformable convolution gives wrong results, what are the most likely bugs?

```
# Common Bug 1: Wrong offset extraction
delta_y = delta[n, k, h_out, w_out]      # WRONG - missing factor of 2
delta_x = delta[n, k + 1, h_out, w_out]   # WRONG - should be 2*k+1

# Correct version:
delta_y = delta[n, ____ * k, h_out, w_out]
delta_x = delta[n, ____ * k + ____, h_out, w_out]

# Common Bug 2: Wrong bilinear interpolation order
# WRONG: Interpolate Y first
v_y = v_00 * (1 - dy) + v_10 * dy
v_final = v_y * (1 - dx) + v_01 * dx

# Correct: Interpolate X first, then Y
v_0 = v_00 * (1 - ____ ) + v_01 * ____   # top edge
v_1 = v_10 * (1 - ____ ) + v_11 * ____   # bottom edge
out = v_0 * (1 - ____ ) + v_1 * ____     # Y interpolation
```

## Scenario 2: CNN Architecture Variations

### Problem Scenario 2: CNN Architecture Variations

If asked to modify the CNN, remember these patterns:

```
# Adding more conv layers - maintain the pattern
class CustomCNNDeep(nn.Module):
    def __init__(self):
        super().__init__()
        # Pattern: start with 3 channels, double each time
        self.conv1 = nn.Conv2d(3, 32, 3, padding=1)
        self.conv2 = nn.Conv2d(32, 32, 3, padding=1) # Same channels
        self.conv3 = nn.Conv2d(32, 64, 3, padding=1) # Double
        self.conv4 = nn.Conv2d(64, 64, 3, padding=1) # Same
        self.conv5 = nn.Conv2d(64, 128, 3, padding=1) # Double

    def forward(self, x):
        # Pattern: Conv-Conv-Pool, Conv-Conv-Pool
        x = F.relu(self.conv1(x))
        x = F.relu(self.conv2(x))
        x = self.pool(x) # After every 2 conv layers

        x = F.relu(self.conv3(x))
        x = F.relu(self.conv4(x))
        x = self.pool(x)

        # Calculate new flatten size: 128 * 8 * 8 = ____
```

## Scenario 3: Validation vs Training Mode

## Problem Scenario 3: Validation vs Training Mode

Critical differences between training and validation:

```
# Training mode
def train_epoch():
    model.____() # Enable dropout and batch norm training mode

    for batch in train_loader:
        optimizer.____() # Clear gradients
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.____() # Compute gradients
        optimizer.____() # Update weights

# Validation mode
def validate():
    model.____() # Disable dropout, batch norm in eval mode

    with torch.____(): # Disable gradient computation
        for batch in val_loader:
            # NO optimizer.zero_grad() here!
            # NO loss.backward() here!
            # NO optimizer.step() here!
            outputs = model(images)
            loss = criterion(outputs, labels)
```

Scenario 4: RNN with Different Sequence Lengths

## Problem Scenario 4: RNN with Different Sequence Lengths

If given a different text, adapt the RNN code:

```
# Original: "Deep Learning"
text = "Deep Learning"
input_seq = text[:-1] # "Deep Learnin"
target_seq = text[1:] # "eep Learning"

# New text: "Hello World"
text = "Hello World"
input_seq = text[____] # "Hello Worl"
target_seq = text[____] # "ello World"

# Vocabulary size changes!
chars = sorted(list(set(text)))
V = len(chars) # This will be different!

# All weight matrices need to be reinitialized with new V
W_xh = torch.randn(H, ____, requires_grad=True) * 0.1
W_hy = torch.randn(____, H, requires_grad=True) * 0.1
b_y = torch.zeros(____, requires_grad=True)
```

Scenario 5: Hyperparameter Grid Search Pattern

## Problem Scenario 5: Hyperparameter Grid Search Pattern

Standard grid search implementation:

```

learning_rates = [0.0001, 0.001]
optimizers = [torch.optim.Adam, torch.optim.SGD]
model_classes = [CustomCNN, CustomCNNwithBN]

best_accuracy = 0
best_params = None

for model_class in model_classes:
    for optimizer_class in optimizers:
        for lr in learning_rates:
            # CRITICAL: Reinitialize model each time
            model = model_class().to(device)

            # Initialize optimizer based on type
            if optimizer_class == torch.optim.SGD:
                optimizer = optimizer_class(
                    model.parameters(),
                    lr=lr,
                    momentum=____,
                    weight_decay=____
                )
            else: # Adam
                optimizer = optimizer_class(
                    model.parameters(),
                    lr=lr,
                    weight_decay=____
                )

            # Train and validate...
            val_acc = train_and_validate(model, optimizer)

            if val_acc > best_accuracy:
                best_accuracy = val_acc
                best_params = (model_class.__name__, optimizer_class.__name__, lr)

```

Scenario 6: Top-K Accuracy Calculation Variations

## Problem Scenario 6: Top-K Accuracy Calculation Variations

Different ways to calculate accuracy:

```

# Top-1 accuracy (most common)
_, predicted = torch.max(outputs, 1)
correct = (predicted == labels).sum().item()
accuracy = correct / labels.size(0) * 100

# Top-K accuracy using topk
_, top_k_pred = outputs.topk(____, dim=1, largest=True, sorted=True)
top_k_correct = top_k_pred.eq(labels.view(-1, 1).expand_as(____)).sum().item()

# Alternative top-K calculation
values, indices = torch.topk(outputs, k=5, dim=1)
correct_mask = indices == labels.unsqueeze(1)
top_k_accuracy = correct_mask.sum().float() / labels.size(0) * 100

```

# EXAM SIMULATION QUESTIONS

Quick Code Writing

## Problem Quick Code Writing

Write these functions from memory in 2 minutes each:

1. Write the bilinear interpolation weight calculation:

```
def calculate_bilinear_weight(p_x, p_y, q_x, q_y):  
    # Your code here - calculate G(p,q)  
    return ____
```

2. Write the RNN hidden state update:

```
def rnn_step(x_t, h_prev, W_xh, W_hh, b_xh, b_hh):  
    # Your code here - compute new hidden state  
    return ____
```

3. Write the CNN forward pass for one block:

```
def cnn_block_forward(x, conv1, conv2, pool):  
    # Your code here - conv-conv-pool pattern  
    return ____
```

4. Write the training step:

```
def training_step(model, optimizer, criterion, images, labels):  
    # Your code here - complete training step  
    return loss
```

5. Write character to one-hot conversion:

```
def char_to_onehot(char, char2idx, vocab_size):  
    # Your code here - convert character to one-hot vector  
    return ____
```

## FINAL MEMORY CHECK

Last Minute Review

## Problem Last Minute Review

Before the exam, quickly verify you remember:

Constants (write from memory):

- CIFAR-100 mean: [\_\_\_\_, \_\_\_\_, \_\_\_\_]
- CIFAR-100 std: [\_\_\_\_, \_\_\_\_, \_\_\_\_]

- Train/val split ratio: \_\_\_\_
- Dropout probability: \_\_\_\_
- Weight decay: \_\_\_\_
- SGD momentum: \_\_\_\_

**Dimensions (write from memory):**

- CNN channels:  $3 \rightarrow \_\_\_ \rightarrow \_\_\_ \rightarrow \_\_\_ \rightarrow \_\_\_$
- CNN spatial:  $32 \rightarrow \_\_\_ \rightarrow \_\_\_ \rightarrow \_\_\_$
- RNN  $W_{xh}$ : (\_\_\_\_, \_\_\_\_)
- RNN  $W_{hh}$ : (\_\_\_\_, \_\_\_\_)
- RNN  $W_{hy}$ : (\_\_\_\_, \_\_\_\_)



# PRACTICE EXAM - NO PEEKING AT ANSWERS!

## Practice 1: Bilinear Interpolation Implementation

### Problem Practice 1: Bilinear Interpolation Implementation

Complete this bilinear interpolation function:

```
def bilinear_interpolate(a_l, q_y, q_x):
    H, W = a_l.shape

    y0 = int(np._____(q_y))
    x0 = int(np._____(q_x))
    y1 = y0 + ___
    x1 = x0 + ___

    def get_pixel_value(img, y, x):
        if 0 <= y < H and 0 <= x < W:
            return img[y, x]
        else:
            return ____

    v_00 = get_pixel_value(a_l, y0, x0)
    v_01 = get_pixel_value(a_l, y0, x1)
    v_10 = get_pixel_value(a_l, y1, x0)
    v_11 = get_pixel_value(a_l, y1, x1)

    dy = q_y - ___
    dx = q_x - ___

    v_0 = v_00 * (1 - ___) + v_01 * ____
    v_1 = v_10 * (1 - ___) + v_11 * ____
    out = v_0 * (1 - ___) + v_1 * ___

    return out
```

## Practice 2: Deformable Conv Offset Extraction

### Problem Practice 2: Deformable Conv Offset Extraction

Fill in the correct offset extraction pattern:

```
for kh in range(K_h):
    for kw in range(K_w):
        k = kh * K_w + kw

        delta_y = delta[n, ___ * k, h_out, w_out]
        delta_x = delta[n, ___ * k + ___, h_out, w_out]
        m_k = mask[n, k, h_out, w_out]

        sample_y = h_start + kh * dilation + ____
        sample_x = w_start + kw * dilation + ____
```

```
# Apply weight and mask
value += weight[c_out, c_in, kh, kw] * ___ * ____
```

### Practice 3: CNN Architecture

## Problem Practice 3: CNN Architecture

Complete the CustomCNN constructor with correct channel progression:

```
class CustomCNN(nn.Module):
    def __init__(self):
        super(CustomCNN, self).__init__()

        self.conv1 = nn.Conv2d(____, ____, kernel_size=3, padding=1)
        self.conv2 = nn.Conv2d(____, ____, kernel_size=3, padding=1)
        self.conv3 = nn.Conv2d(____, ____, kernel_size=3, padding=1)
        self.conv4 = nn.Conv2d(____, ____, kernel_size=3, padding=1)

        self.pool = nn._____(2, 2)

        self.fc1 = nn.Linear(___ * ___ * ____, 512)
        self.fc2 = nn.Linear(____, ____)
        self.fc3 = nn.Linear(____, ____)

        self.dropout = nn.Dropout(____)
```

### Practice 4: CNN Forward Pass

## Problem Practice 4: CNN Forward Pass

Complete the forward pass with correct activations:

```
def forward(self, x):
    x = F.____(self.conv1(x))
    x = F.____(self.conv2(x))
    x = self.pool(x) # 32x32 -> ____

    x = F.____(self.conv3(x))
    x = self.pool(x) # 16x16 -> ____

    x = F.____(self.conv4(x))
    x = self.pool(x) # 8x8 -> ____

    x = x.view(x.____(), ____)

    x = F.____(self.fc1(x))
    x = self._____(x)
```

```

x = F.____(self.fc2(x))
x = self._____(x)
x = self.fc3(x)

return x

```

Practice 5: Training Loop Pattern

## Problem Practice 5: Training Loop Pattern

Complete the sacred training loop:

```

def train_step(model, optimizer, criterion, images, labels, device):
    images = images.to(____)
    labels = labels.to(____)

    optimizer.____()
    outputs = model(_____)
    loss = criterion(outputs, labels)
    loss._____(____)
    optimizer.____()

    return loss.item()

```

Practice 6: Top-K Accuracy

## Problem Practice 6: Top-K Accuracy

Complete the top-1 and top-5 accuracy calculation:

```

# Top-1 accuracy
_, top1_pred = outputs.topk(____, dim=1, largest=____, sorted=____)
top1_correct = top1_pred.eq(labels.view(____, ____)).sum().item()

# Top-5 accuracy
_, top5_pred = outputs.topk(____, dim=1, largest=____, sorted=____)
top5_correct = top5_pred.eq(labels.view(-1, 1).____(____)).sum().item()

```

Practice 7: RNN Weight Initialization

## Problem Practice 7: RNN Weight Initialization

Initialize RNN parameters with correct dimensions:

```

V = len(chars) # Vocabulary size
H = 16         # Hidden size

W_xh = torch.randn(____, ____, requires_grad=True) * 0.1
W_hh = torch.randn(____, ____, requires_grad=True) * 0.1
b_xh = torch.zeros(____, requires_grad=True)
b_hh = torch.zeros(____, requires_grad=True)
W_hy = torch.randn(____, ____, requires_grad=True) * 0.1
b_y = torch.zeros(____, requires_grad=True)

```

Practice 8: RNN Forward Equations

## Problem Practice 8: RNN Forward Equations

Complete the RNN forward pass:

```

for t in range(seq_len):
    x_t = inputs[t]

    # Hidden state update
    h = torch.tanh(_____ @ x_t + _____ + _____ @ h + _____)

    # Output logits
    s_t = _____ @ h + _____

    logits_list.append(s_t)

```

Practice 9: Character Processing

## Problem Practice 9: Character Processing

Complete the character-level preprocessing:

```

text = "Deep Learning"

chars = sorted(list(set(_____)))
char2idx = {ch: i for i, ch in enumerate(_____)}
idx2char = {i: ch for i, ch in enumerate(_____)}

input_seq = text[_____]
target_seq = text[_____]

def one_hot(idx, size):
    vec = torch.zeros(size)
    vec[_____] = 1.0
    return vec

```

```
inputs = [one_hot(char2idx[ch], V) for ch in ____]
targets = torch.tensor([char2idx[ch] for ch in ____], dtype=torch.long)
```

## Practice 10: Gradient Computation

### Problem Practice 10: Gradient Computation

Complete the explicit gradient computation:

```
grad_W_xh = torch.autograd.grad(loss_manual, W_xh, retain_graph=____)[0]
grad_W_hh = torch.autograd.grad(loss_manual, W_hh, retain_graph=____)[0]
grad_b_xh = torch.autograd.grad(loss_manual, b_xh, retain_graph=____)[0]
grad_b_hh = torch.autograd.grad(loss_manual, b_hh, retain_graph=____)[0]
grad_W_hy = torch.autograd.grad(loss_manual, W_hy, retain_graph=____)[0]
grad_b_y = torch.autograd.grad(loss_manual, b_y, retain_graph=____)[0]
```

## Practice 11: BatchNorm Integration

### Problem Practice 11: BatchNorm Integration

Add BatchNorm to CNN architecture:

```
class CustomCNNwithBN(nn.Module):
    def __init__(self):
        super().__init__()

        self.conv1 = nn.Conv2d(3, 32, kernel_size=3, padding=1)
        self.bn1 = nn.BatchNorm2d(____)

        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, padding=1)
        self.bn2 = nn.BatchNorm2d(____)

    def forward(self, x):
        x = F.relu(self.____(self.conv1(x)))
        x = F.relu(self.____(self.conv2(x)))
```

## Practice 12: CIFAR-100 Constants

### Problem Practice 12: CIFAR-100 Constants

Fill in the CIFAR-100 normalization values:

```
transform = transforms.Compose([
```

```

        transforms.ToTensor(),
        transforms.Normalize(mean=[____, ____, ____],
                               std=[____, ____, ____])
    ])

```

### Practice 13: Data Splitting

## Problem Practice 13: Data Splitting

Complete the 80/20 train/validation split:

```

full_train_set = CIFAR100(root='./data', train=True, download=True)

train_size = int(____ * len(full_train_set))
val_size = len(full_train_set) - ____

train_set, val_set = random_split(full_train_set, [____, ____])

```

### Practice 14: Loss and Optimizer Setup

## Problem Practice 14: Loss and Optimizer Setup

Set up loss function and optimizer for CIFAR-100:

```

loss_function = nn.______()

optimizer = optim.SGD(model.parameters(),
                       lr=____,
                       momentum=____,
                       weight_decay=____)

```

### Practice 15: Quick Recall

## Problem Practice 15: Quick Recall

Answer these without looking:

1. What function rounds down to nearest integer? \_\_\_\_\_
2. In PyTorch deformable conv, Y offset comes at index \_\_\_\_ and X offset at \_\_\_\_
3. The bilinear interpolation weight formula is:  $G(p,q) = \text{_____}$

4. CNN channel progression in our model:  $3 \rightarrow \_\_\_ \rightarrow \_\_\_ \rightarrow \_\_\_ \rightarrow \_\_\_$
5. RNN hidden state uses which activation function? \_\_\_\_\_
6. In training loop, what comes after `loss.backward()`? \_\_\_\_\_
7. For validation, you wrap the loop with `torch._____()`:
8. The flatten operation uses `x._____(x._____(0), ____)`:
9. Top-5 accuracy uses `outputs._____(____, dim=1)`:
10. RNN weight `W_hy` has dimensions `(____, ____)`:

## CHALLENGE PROBLEMS

Challenge 1: Complete Function Implementation

### **Problem Challenge 1: Complete Function Implementation**

Write the complete bilinear interpolation function from scratch:

Challenge 2: Complete RNN Step

### **Problem Challenge 2: Complete RNN Step**

Write a complete RNN forward step function:

Challenge 3: Training vs Validation

### **Problem Challenge 3: Training vs Validation**

Write the key differences between training and validation loops:



#### Challenge 4: Debugging Scenario

### Problem Challenge 4: Debugging Scenario

Your deformable convolution outputs don't match PyTorch. List 5 most likely bugs:

1. \_\_\_\_\_
2. \_\_\_\_\_
3. \_\_\_\_\_
4. \_\_\_\_\_
5. \_\_\_\_\_

#### Challenge 5: Architecture Design

### Problem Challenge 5: Architecture Design

Design a deeper CNN with 6 conv layers following the same pattern:

## ANSWER BANK

Practice 1 Answers:

### Problem Practice 1 Answers:

floor, floor, 1, 1, 0.0, y0, x0, dx, dx, dx, dx, dy, dy

Practice 2 Answers:

### Problem Practice 2 Answers:

2, 2, 1, delta\_y, delta\_x, m\_k, interpolated

Practice 3 Answers:

### Problem Practice 3 Answers:

3, 32, 32, 64, 64, 128, 128, 256, MaxPool2d, 256, 4, 4, 512, 256, 256, 100, 0.5

Practice 4 Answers:

### Problem Practice 4 Answers:

relu, relu, 16x16, relu, 8x8, relu, 4x4, size, 0, -1, relu, dropout, relu, dropout

Practice 5 Answers:

### Problem Practice 5 Answers:

device, device, zero\_grad, images, backward, step

Practice 6 Answers:

### Problem Practice 6 Answers:

1, True, True, -1, 1, 5, True, True, expand\_as(top5\_pred)

Practice 7 Answers:

### Problem Practice 7 Answers:

H, V, H, H, H, H, V, H, V

Practice 8 Answers:

### Problem Practice 8 Answers:

W\_xh, b\_xh, W\_hh, b\_hh, W\_hy, b\_y

Practice 9 Answers:

### Problem Practice 9 Answers:

text, chars, chars, [:-1], [1:], idx, input\_seq, target\_seq

Practice 10 Answers:

### Problem Practice 10 Answers:

True, True, True, True, True, True

Practice 11 Answers:

### Problem Practice 11 Answers:

32, 64, bn1, bn2

Practice 12 Answers:

### Problem Practice 12 Answers:

0.5071, 0.4867, 0.4408, 0.2675, 0.2565, 0.2761

Practice 13 Answers:

### Problem Practice 13 Answers:

0.8, train\_size, train\_size, val\_size

Practice 14 Answers:

### Problem Practice 14 Answers:

CrossEntropyLoss, 0.01, 0.9, 5e-4

Practice 15 Answers:

### Problem Practice 15 Answers:

1. np.floor 2. 2\*k, 2\*k+1 3.  $(1 - p_x - q_x) * (1 - p_y - q_y)$  4. 32, 64, 128, 256 5. tanh 6. optimizer.step()  
7. no\_grad 8. view, size, -1 9. topk, 5 10. V, H

Challenge Problem Answers:

### Problem Challenge Problem Answers:

**Challenge 1:** Complete bilinear function with all bounds checking, corner calculations, and interpolation steps.

**Challenge 2:** RNN step with tanh activation: `h = torch.tanh(W_xh @ x + b_xh + W_hh @ h_prev + b_hh)`

**Challenge 3:** Training: `model.train()`, optimizer steps, gradients enabled. Validation: `model.eval()`, `torch.no_grad()`, no optimizer steps.

**Challenge 4:** 1. Wrong offset indexing (2\*k), 2. Wrong bilinear order (X first), 3. Missing zero padding, 4. Wrong sampling position calculation, 5. Incorrect weight/mask application.

**Challenge 5:** Conv layers: 3→32→32→64→64→128→128, with pooling after every 2 conv layers.

# CENG403 - Spring 2025: Homework set THE-2 Study Guide

Your Name

## TASK 1: DEFORMABLE CNN MEMORIZATION GUIDE

.1

### Problem 1.1

What are the 4 corner positions for bilinear interpolation given fractional position  $q = (q_x, q_y)$ ? **Pattern to Remember:** Floor-Floor, Ceil-Floor, Floor-Ceil, Ceil-Ceil

$$p_{lt} = (\lfloor q_x \rfloor, \lfloor q_y \rfloor) \quad (\text{left top}) \quad (1)$$

$$p_{rt} = (\lceil q_x \rceil, \lfloor q_y \rfloor) \quad (\text{right top}) \quad (2)$$

$$p_{lb} = (\lfloor q_x \rfloor, \lceil q_y \rceil) \quad (\text{left bottom}) \quad (3)$$

$$p_{rb} = (\lceil q_x \rceil, \lceil q_y \rceil) \quad (\text{right bottom}) \quad (4)$$

.2

### Problem 1.2

Complete the bilinear interpolation weight formula:

$$G(p, q) = (1 - |p_x - q_x|) \cdot (1 - |p_y - q_y|)$$

.3

### Problem 1.3

Fill in the missing code for bilinear interpolation bounds checking:

```
def get_pixel_value(img, y, x):  
    if 0 <= y < H and 0 <= x < W:  
        return img[y, x]  
    else:  
        return ___ # What goes here?
```

**Answer:** 0.0 (zero padding for out-of-bounds)

.4

### Problem 1.4

What is the correct order for bilinear interpolation calculation? **Memory Pattern:** "First X, then Y"

1. Get 4 corner values:  $v_{00}, v_{01}, v_{10}, v_{11}$
2. Calculate fractional parts:  $dx = q_x - x_0, dy = q_y - y_0$
3. Interpolate along X:  $v_0 = v_{00}(1 - dx) + v_{01} \cdot dx$
4. Interpolate along X:  $v_1 = v_{10}(1 - dx) + v_{11} \cdot dx$
5. Interpolate along Y:  $out = v_0(1 - dy) + v_1 \cdot dy$

.5

## Problem 1.5

In deformable convolution, how do you extract the y and x offsets from the delta tensor? **Critical Pattern - PyTorch stores Y first, then X:**

```
delta_y = delta[n, 2 * k, h_out, w_out]    # y offset
delta_x = delta[n, 2 * k + 1, h_out, w_out] # x offset
```

.6

## Problem 1.6

What is the deformable convolution sampling position formula?

```
sample_y = h_start + kh * dilation + _____
sample_x = w_start + kw * dilation + _____
```

**Answer:** delta\_y and delta\_x

## TASK 2: CNN PYTORCH MEMORIZATION GUIDE

.1

### Problem 2.1

What are the CIFAR-100 normalization values you must memorize? **Critical Constants:**

```
mean=[0.5071, 0.4867, 0.4408] # CIFAR100 mean
std=[0.2675, 0.2565, 0.2761]  # CIFAR100 std
```

.2

### Problem 2.2

Complete the data augmentation transforms for training:

```
transform_train = transforms.Compose([
    transforms._____(32, padding=4), # What goes here?
    transforms._____,               # What goes here?
    transforms.ToTensor(),
    transforms.Normalize(mean=[...], std=[...])
])
```

**Answer:** RandomCrop and RandomHorizontalFlip

.3

### Problem 2.3

How do you split CIFAR-100 training data into 80/20 train/validation? **Pattern to Remember:**

```
train_size = int(0.8 * len(full_train_set))
val_size = len(full_train_set) - train_size
train_set, val_set = random_split(full_train_set, [train_size, val_size])
```

.4

### Problem 2.4

What is the CNN architecture pattern for the CustomCNN class? **Layer Sequence Pattern:**

1. Conv2d(3, 32) → Conv2d(32, 64) → MaxPool2d
2. Conv2d(64, 128) → MaxPool2d
3. Conv2d(128, 256) → MaxPool2d
4. Flatten → FC(256\*4\*4, 512) → FC(512, 256) → FC(256, 100)

.5

### Problem 2.5

Complete the forward pass activation pattern:

```

x = F.relu(self.conv1(x))
x = F.relu(self.conv2(x))
x = self.pool(x) # 32x32 -> 16x16
x = F.relu(self.conv3(x))
x = self.pool(x) # 16x16 -> 8x8
x = F.relu(self.conv4(x))
x = self.pool(x) # 8x8 -> 4x4
x = x.view(x.size(0), -1) # Flatten
x = F.relu(self.fc1(x))
x = self.dropout(x)
x = F.relu(self.fc2(x))
x = self.dropout(x)
x = self.fc3(x) # No activation on final layer!

```

.6

## Problem 2.6

What loss function and optimizer setup is standard for CIFAR-100?

```

loss_function = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(),
                      lr=0.01,
                      momentum=0.9,
                      weight_decay=5e-4)

```

.7

## Problem 2.7

Complete the top-1 and top-5 accuracy calculation:

```

# Top-1 accuracy
_, top1_pred = outputs.topk(1, dim=1, largest=True, sorted=True)
top1_correct = top1_pred.eq(labels.view(-1, 1)).sum().item()

# Top-5 accuracy
_, top5_pred = outputs.topk(5, dim=1, largest=True, sorted=True)
top5_correct = top5_pred.eq(labels.view(-1, 1).____(____)).sum().item()

```

Answer: `expand_as(top5_pred)`

.8

## Problem 2.8

What is the training loop structure pattern? **Memory Pattern - "Zero, Forward, Backward, Step"**:

```

optimizer.zero_grad() # Clear gradients
outputs = model(images) # Forward pass
loss = loss_function(outputs, labels) # Compute loss
loss.backward() # Backward pass
optimizer.step() # Update weights

```

.9

## Problem 2.9

How do you add BatchNorm2d to the CNN architecture? **Pattern - After each Conv2d:**

```
self.conv1 = nn.Conv2d(3, 32, kernel_size=3, padding=1)
self.bn1 = nn.BatchNorm2d(32) # Same number as conv output
self.conv2 = nn.Conv2d(32, 64, kernel_size=3, padding=1)
self.bn2 = nn.BatchNorm2d(64) # Same number as conv output
```



## TASK 3: RNN MEMORIZATION GUIDE

.1

### Problem 3.1

How do you create character vocabulary and mappings? **Standard Pattern:**

```
chars = sorted(list(set(text)))
char2idx = {ch: i for i, ch in enumerate(chars)}
idx2char = {i: ch for i, ch in enumerate(chars)}
```

.2

### Problem 3.2

How do you create input and target sequences for character prediction?

```
input_seq = text[:-1] # All except last
target_seq = text[1:] # All except first
```

.3

### Problem 3.3

Complete the one-hot encoding function:

```
def one_hot(idx, size):
    vec = torch.zeros(size)
    vec[idx] = 1.0
    return vec
```

**Answer:** idx

.4

### Problem 3.4

What are the RNN weight matrix dimensions? **Dimension Memory Pattern:**

```
W_xh = torch.randn(H, V, requires_grad=True) * 0.1 # (H, V)
W_hh = torch.randn(H, H, requires_grad=True) * 0.1 # (H, H)
b_xh = torch.zeros(H, requires_grad=True)           # (H,)
b_hh = torch.zeros(H, requires_grad=True)           # (H,)
W_hy = torch.randn(V, H, requires_grad=True) * 0.1 # (V, H)
b_y = torch.zeros(V, requires_grad=True)            # (V,)
```

.5

### Problem 3.5

Complete the RNN forward pass equations:

```
# Hidden state update
h = torch.tanh(W_xh @ x_t + b_xh + W_hh @ h + b_hh)
```

```
# Output logits
s_t = ----- @ h + -----
```

**Answer:**  $W_{hy}$  and  $b_y$

.6

## Problem 3.6

How do you compute gradients explicitly with `torch.autograd`?

```
grad_W_xh = torch.autograd.grad(loss_manual, W_xh, retain_graph=True)[0]
grad_W_hh = torch.autograd.grad(loss_manual, W_hh, retain_graph=True)[0]
grad_b_xh = torch.autograd.grad(loss_manual, b_xh, retain_graph=True)[0]
grad_b_hh = torch.autograd.grad(loss_manual, b_hh, retain_graph=True)[0]
grad_W_hy = torch.autograd.grad(loss_manual, W_hy, retain_graph=True)[0]
grad_b_y = torch.autograd.grad(loss_manual, b_y, retain_graph=True)[0]
```

# CRITICAL MEMORIZATION PATTERNS

Patterns

## Problem Patterns

What are the key patterns you must memorize?

### 1. Bilinear Interpolation Pattern:

- Get 4 corners (floor/ceil combinations)
- Interpolate X first, then Y
- Use fractional parts:  $dx = q_x - x_0$ ,  $dy = q_y - y_0$

### 2. CNN Architecture Pattern:

- Conv-Conv-Pool, Conv-Pool, Conv-Pool structure
- Channel progression:  $3 \rightarrow 32 \rightarrow 64 \rightarrow 128 \rightarrow 256$
- Spatial reduction:  $32 \rightarrow 16 \rightarrow 8 \rightarrow 4$
- FC layers:  $4096 \rightarrow 512 \rightarrow 256 \rightarrow 100$

### 3. Training Loop Pattern:

- "Zero-Forward-Backward-Step"
- Always move tensors to device
- Use `torch.no_grad()` for validation

### 4. RNN Equations Pattern:

- $h_t = \tanh(W_{xh}x_t + b_{xh} + W_{hh}h_{t-1} + b_{hh})$
- $s_t = W_{hy}h_t + b_y$
- Always remember matrix dimensions

### 5. Data Preprocessing Patterns:

- CIFAR-100: `mean=[0.5071, 0.4867, 0.4408]`, `std=[0.2675, 0.2565, 0.2761]`
- 80/20 split: `train_size = int(0.8 * len(dataset))`
- Character sequences: `input = text[:-1]`, `target = text[1:]`

# COMMON MISTAKES TO AVOID

Mistakes

## Problem Mistakes

What are the most common implementation mistakes?

### 1. Deformable CNN:

- Forget PyTorch stores Y offset first, then X offset
- Wrong bilinear interpolation order (do X first, then Y)
- Forget zero padding for out-of-bounds pixels

## 2. CNN:

- Forget to move tensors to device
- Wrong flatten calculation: `x.view(x.size(0), -1)`
- Forget activation on hidden layers, add activation on output layer

## 3. RNN:

- Wrong weight matrix dimensions
- Forget `requires_grad=True` for parameters
- Use `retain_graph=True` for multiple gradient computations

# RAPID-FIRE ACTIVE RECALL QUIZ

Speed Round 1: Fill the Blanks

## Problem Speed Round 1: Fill the Blanks

Complete these critical code snippets:

**Q1:** Bilinear interpolation corners:

```
y0 = int(np._____(q_y))
x0 = int(np._____(q_x))
y1 = y0 + ____
x1 = x0 + ____
```

**Q2:** Deformable conv offset extraction:

```
delta_y = delta[n, ____ * k, h_out, w_out]
delta_x = delta[n, ____ * k + ____, h_out, w_out]
```

**Q3:** CNN flatten operation:

```
x = x.view(x.____(__), ____)
```

**Q4:** Top-5 accuracy calculation:

```
_, top5_pred = outputs.topk(____, dim=1, largest=____, sorted=____)
top5_correct = top5_pred.eq(labels.view(-1, 1)._____(____)).sum().item()
```

**Q5:** RNN hidden state update:

```
h = torch.tanh(____ @ x_t + _____ + _____ @ h + _____)
```

Speed Round 2: True/False

## Problem Speed Round 2: True/False

Mark T/F for these statements:

1. In bilinear interpolation, you interpolate Y direction first, then X direction. **[T/F]**
2. PyTorch stores Y offset before X offset in deformable convolution. **[T/F]**
3. CIFAR-100 has 100 classes, so the final FC layer outputs 100 values. **[T/F]**
4. You should apply ReLU activation to the final output layer in classification. **[T/F]**
5. In RNN,  $W_{xh}$  has dimensions (V, H). **[T/F]**
6. For validation, you need to call `optimizer.zero_grad()`. **[T/F]**
7. BatchNorm2d should be applied before the activation function. **[T/F]**
8. The input sequence for RNN is `text[1:]` and target is `text[:-1]`. **[T/F]**

### Speed Round 3: Memory Palace

## Problem Speed Round 3: Memory Palace

Associate these concepts with memorable phrases:

**Bilinear Interpolation:** "Four corners, X then Y, fractional magic"

- 4 corners: lt, rt, lb, rb
- X interpolation: top\_edge, bottom\_edge
- Y interpolation: final result

**CNN Architecture:** "3 to 32, double-double-double, then shrink to 100"

- Channels:  $3 \rightarrow 32 \rightarrow 64 \rightarrow 128 \rightarrow 256$
- Spatial:  $32 \rightarrow 16 \rightarrow 8 \rightarrow 4$
- FC:  $4096 \rightarrow 512 \rightarrow 256 \rightarrow 100$

**Training Loop:** "Zero-Forward-Backward-Step dance"

- `optimizer.zero_grad()`
- `outputs = model(images)`
- `loss.backward()`
- `optimizer.step()`

**RNN Weights:** "Input-Hidden-Hidden, Hidden-Hidden-Hidden, Hidden-Vocab-Vocab"

- $W_{xh}$ : (H, V) - maps input to hidden
- $W_{hh}$ : (H, H) - maps previous hidden to current hidden
- $W_{hy}$ : (V, H) - maps hidden to output vocabulary

## LAST-MINUTE CHECKLIST

### Pre-Exam Checklist

## Problem Pre-Exam Checklist

Before the exam, ensure you can write from memory:

**Critical Constants:**

- CIFAR-100 mean: [0.5071, 0.4867, 0.4408]
- CIFAR-100 std: [0.2675, 0.2565, 0.2761]
- Train/val split:  $0.8 * \text{len}(\text{dataset})$

**Key Formulas:**

- Bilinear weight:  $(1 - |p_x - q_x|) \cdot (1 - |p_y - q_y|)$
- RNN hidden:  $h_t = \tanh(W_{xh}x_t + b_{xh} + W_{hh}h_{t-1} + b_{hh})$
- RNN output:  $s_t = W_{hy}h_t + b_y$

#### Critical Code Patterns:

- Device transfer: `tensor.to(device)`
- Gradient computation: `torch.autograd.grad(loss, param, retain_graph=True)[0]`
- Top-k accuracy: `outputs.topk(k, dim=1, largest=True, sorted=True)`
- One-hot encoding: `vec[idx] = 1.0`

#### Architecture Patterns:

- CNN: Conv→BN→ReLU→Pool pattern
- RNN: Input→Hidden→Output with recurrence
- Training: Zero→Forward→Backward→Step

# CODING BLOCKS MEMORIZATION

Code Block 1: Bilinear Interpolation Core

## Problem Code Block 1: Bilinear Interpolation Core

Complete this bilinear interpolation function - focus on the mathematical pattern:

```
def bilinear_interpolate(a_l, q_y, q_x):
    H, W = a_l.shape

    # Step 1: Get integer positions
    y0 = int(np._____(q_y))
    x0 = int(np._____(q_x))
    y1 = y0 + ___
    x1 = x0 + ___

    # Step 2: Get values with bounds checking
    def get_pixel_value(img, y, x):
        if 0 <= y < H and 0 <= x < W:
            return img[y, x]
        else:
            return ____ # Out of bounds value

    # Step 3: Get four corner values
    v_00 = get_pixel_value(a_l, y0, x0) # ____-____
    v_01 = get_pixel_value(a_l, y0, x1) # ____-____
    v_10 = get_pixel_value(a_l, y1, x0) # ____-____
    v_11 = get_pixel_value(a_l, y1, x1) # ____-____

    # Step 4: Calculate fractional parts
    dy = q_y - ___
    dx = q_x - ___

    # Step 5: Interpolate X first, then Y
    v_0 = v_00 * (1 - ___) + v_01 * ___ # top edge
    v_1 = v_10 * (1 - ___) + v_11 * ___ # bottom edge
    out = v_0 * (1 - ___) + v_1 * ___ # final Y interpolation

    return out
```

Code Block 2: Deformable Conv Key Loop

## Problem Code Block 2: Deformable Conv Key Loop

This is the heart of deformable convolution - memorize the offset extraction pattern:

```
for n in range(N): # batch
    for c_out in range(C_out): # output channels
        for h_out in range(H_out): # height
            for w_out in range(W_out): # width
                h_start = h_out * ____
                w_start = w_out * ____
                value = 0.0

                for kh in range(K_h):
                    for kw in range(K_w):
```



```

k = kh * K_w + kw

# CRITICAL: PyTorch offset order
delta_y = delta[n, ___ * k, h_out, w_out]
delta_x = delta[n, ___ * k + ___, h_out, w_out]
m_k = mask[n, k, h_out, w_out]

# Sampling position
sample_y = h_start + kh * dilation + ____
sample_x = w_start + kw * dilation + ____

for c_in in range(C_in):
    interpolated = bilinear_interpolate(
        a_l[n, c_in, :, :], sample_y, sample_x
    )
    value += weight[c_out, c_in, kh, kw] * ___ * ____

out[n, c_out, h_out, w_out] = value

```

Code Block 3: CNN Architecture Constructor

## Problem Code Block 3: CNN Architecture Constructor

Memorize the channel progression and layer naming pattern:

```

class CustomCNN(nn.Module):
    def __init__(self, norm_layer=None):
        super(CustomCNN, self).__init__()

        # Conv layers - memorize the channel progression
        self.conv1 = nn.Conv2d(____, ___, kernel_size=3, padding=1)
        self.conv2 = nn.Conv2d(____, ___, kernel_size=3, padding=1)
        self.conv3 = nn.Conv2d(____, ___, kernel_size=3, padding=1)
        self.conv4 = nn.Conv2d(____, ___, kernel_size=3, padding=1)

        # Pooling layer
        self.pool = nn._____(2, 2)

        # FC layers - calculate the input size
        self.fc1 = nn.Linear(___ * ___ * ___, 512)
        self.fc2 = nn.Linear(____, ___)
        self.fc3 = nn.Linear(____, ___) # CIFAR-100 classes

        self.dropout = nn.Dropout(____)

```

Code Block 4: CNN Forward Pass Pattern

## Problem Code Block 4: CNN Forward Pass Pattern

Memorize the activation and pooling pattern:

```

def forward(self, x):
    # Block 1: Conv-Conv-Pool
    x = F.____(self.conv1(x))
    x = F.____(self.conv2(x))
    x = self.pool(x) # 32x32 -> ____

```

```

# Block 2: Conv-Pool
x = F.____(self.conv3(x))
x = self.pool(x) # 16x16 -> ____

# Block 3: Conv-Pool
x = F.____(self.conv4(x))
x = self.pool(x) # 8x8 -> ____

# Flatten
x = x.view(x.____(__), __)

# FC layers with dropout
x = F.____(self.fc1(x))
x = self.____(x)
x = F.____(self.fc2(x))
x = self.____(x)
x = self.fc3(x) # No activation here!

return x

```

Code Block 5: Training Loop Core

## Problem Code Block 5: Training Loop Core

The sacred training loop pattern - memorize the order:

```

def train(model, train_loader, optimizer, loss_function, device):
    model.____() # Set to training mode

    for batch in train_loader:
        images, labels = batch
        images = images.to(____)
        labels = labels.to(____)

        # The sacred four steps:
        optimizer.____() # Step 1: Clear gradients
        outputs = model(____) # Step 2: Forward pass
        loss = loss_function(outputs, labels) # Step 3: Compute loss
        loss.____() # Step 4: Backward pass
        optimizer.____() # Step 5: Update weights

        # Accuracy calculation
        _, top1_pred = outputs.topk(____, dim=1, largest=True, sorted=True)
        top1_correct = top1_pred.eq(labels.view(____, ____)).sum().item()

```

Code Block 6: BatchNorm CNN Constructor

## Problem Code Block 6: BatchNorm CNN Constructor

Pattern for adding BatchNorm after each conv layer:

```

class CustomCNNwithBN(nn.Module):
    def __init__(self, norm_layer=None):
        super(CustomCNNwithBN, self).__init__()

        self.conv1 = nn.Conv2d(3, 32, kernel_size=3, padding=1)
        self.bn1 = nn.BatchNorm2d(____) # Same as conv1 output

```

```

self.conv2 = nn.Conv2d(32, 64, kernel_size=3, padding=1)
self.bn2 = nn.BatchNorm2d(____) # Same as conv2 output

self.conv3 = nn.Conv2d(64, 128, kernel_size=3, padding=1)
self.bn3 = nn.BatchNorm2d(____) # Same as conv3 output

self.conv4 = nn.Conv2d(128, 256, kernel_size=3, padding=1)
self.bn4 = nn.BatchNorm2d(____) # Same as conv4 output

def forward(self, x):
    # Pattern: Conv -> BN -> ReLU
    x = F.relu(self.bn1(self.conv1(x)))
    x = F.relu(self.bn2(self.conv2(x)))
    x = self.pool(x)

    x = F.relu(self.bn3(self.conv3(x)))
    x = self.pool(x)

    x = F.relu(self.bn4(self.conv4(x)))
    x = self.pool(x)
    # ... rest of forward pass

```

Code Block 7: RNN Parameter Initialization

## Problem Code Block 7: RNN Parameter Initialization

Memorize the weight dimensions and initialization pattern:

```

# RNN parameters - memorize the dimensions!
W_xh = torch.randn(____, ____, requires_grad=True) * 0.1 # Input to hidden
W_hh = torch.randn(____, ____, requires_grad=True) * 0.1 # Hidden to hidden
b_xh = torch.zeros(____, requires_grad=True)             # Hidden bias 1
b_hh = torch.zeros(____, requires_grad=True)             # Hidden bias 2
W_hy = torch.randn(____, ____, requires_grad=True) * 0.1 # Hidden to output
b_y = torch.zeros(____, requires_grad=True)              # Output bias

# Remember: V = vocab size, H = hidden size

```

Code Block 8: RNN Forward Pass

## Problem Code Block 8: RNN Forward Pass

The RNN equations in code form:

```

logits_list = []
h = torch.zeros(H)

for t in range(seq_len):
    x_t = inputs[t]

    # Hidden state update equation
    h = torch.tanh(____ @ x_t + ____ + ____ @ h + ____ )

    # Output logits equation
    s_t = ____ @ h + ____

```

```

logits_list.append(s_t)

logits = torch.stack(logits_list)
log_probs = F.log_softmax(logits, dim=1)
loss_manual = F.nll_loss(log_probs, targets)

```

Code Block 9: Gradient Computation

## Problem Code Block 9: Gradient Computation

Pattern for explicit gradient computation:

```

# Compute gradients explicitly
grad_W_xh = torch.autograd.grad(loss_manual, W_xh, retain_graph=____)[0]
grad_W_hh = torch.autograd.grad(loss_manual, W_hh, retain_graph=____)[0]
grad_b_xh = torch.autograd.grad(loss_manual, b_xh, retain_graph=____)[0]
grad_b_hh = torch.autograd.grad(loss_manual, b_hh, retain_graph=____)[0]
grad_W_hy = torch.autograd.grad(loss_manual, W_hy, retain_graph=____)[0]
grad_b_y = torch.autograd.grad(loss_manual, b_y, retain_graph=____)[0]

# Why retain_graph=True? Because we compute multiple gradients from same loss

```

Code Block 10: Character Processing

## Problem Code Block 10: Character Processing

Standard pattern for character-level RNN preprocessing:

```

text = "Deep Learning"

# Step 1: Create vocabulary
chars = sorted(list(set(____)))
char2idx = {ch: i for i, ch in enumerate(____)}
idx2char = {i: ch for i, ch in enumerate(____)}

# Step 2: Create sequences
input_seq = text[____] # All except last
target_seq = text[____] # All except first

# Step 3: Convert to tensors
inputs = [one_hot(char2idx[ch], V) for ch in ____]
targets = torch.tensor([char2idx[ch] for ch in ____], dtype=torch.long)

def one_hot(idx, size):
    vec = torch.zeros(size)
    vec[____] = 1.0
    return vec

```

# MEMORIZATION MNEMONICS

Memory Aids

## Problem Memory Aids

Use these phrases to remember key patterns:

**"Floor-Ceil-Four-Corners":** Bilinear interpolation corners

- lt: floor-floor, rt: ceil-floor, lb: floor-ceil, rb: ceil-ceil

**"Y-before-X-in-PyTorch":** Deformable conv offset ordering

- $\text{delta\_y} = \text{delta}[n, 2*k, \text{h\_out}, \text{w\_out}]$
- $\text{delta\_x} = \text{delta}[n, 2*k+1, \text{h\_out}, \text{w\_out}]$

**"3-32-64-128-256":** CNN channel progression

- Each layer doubles the channels (except first)

**"32-16-8-4":** Spatial dimension reduction

- Each `MaxPool2d(2,2)` halves the spatial dimensions

**"Zero-Forward-Backward-Step":** Training loop mantra

- Never forget the order!

**"Input-Hidden-Hidden":** RNN weight dimensions

- $W_{xh}$ : (H,V),  $W_{hh}$ : (H,H),  $W_{hy}$ : (V,H)

**"Tanh-Hidden-Linear-Output":** RNN computation flow

- Hidden uses tanh, output is linear

## ANSWERS TO CODING BLOCKS

**Code Block 1:** floor, floor, 1, 1, 0.0, top-left, top-right, bottom-left, bottom-right, y0, x0, dx, dx, dx, dx, dy, dy

**Code Block 2:** stride, stride, 2, 2, 1, delta\_y, delta\_x, m\_k, interpolated

**Code Block 3:** 3, 32, 32, 64, 64, 128, 128, 256, `MaxPool2d`, 256, 4, 4, 512, 256, 256, 100, 0.5

**Code Block 4:** relu, relu, 16x16, relu, 8x8, relu, 4x4, `size(0)`, -1, relu, dropout, relu, dropout

**Code Block 5:** train, device, device, zero\_grad, images, backward, step, 1, -1, 1

**Code Block 6:** 32, 64, 128, 256

**Code Block 7:** H, V, H, H, H, H, V, H, V

**Code Block 8:** W\_xh, b\_xh, W\_hh, b\_hh, W\_hy, b\_y

**Code Block 9:** True, True, True, True, True, True

# ADVANCED CODING SCENARIOS

## Scenario 1: Debugging Deformable Conv

### Problem Scenario 1: Debugging Deformable Conv

If your deformable convolution gives wrong results, what are the most likely bugs?

```
# Common Bug 1: Wrong offset extraction
delta_y = delta[n, k, h_out, w_out]      # WRONG - missing factor of 2
delta_x = delta[n, k + 1, h_out, w_out]   # WRONG - should be 2*k+1

# Correct version:
delta_y = delta[n, ____ * k, h_out, w_out]
delta_x = delta[n, ____ * k + ____, h_out, w_out]

# Common Bug 2: Wrong bilinear interpolation order
# WRONG: Interpolate Y first
v_y = v_00 * (1 - dy) + v_10 * dy
v_final = v_y * (1 - dx) + v_01 * dx

# Correct: Interpolate X first, then Y
v_0 = v_00 * (1 - ____ ) + v_01 * ____   # top edge
v_1 = v_10 * (1 - ____ ) + v_11 * ____   # bottom edge
out = v_0 * (1 - ____ ) + v_1 * ____     # Y interpolation
```

## Scenario 2: CNN Architecture Variations

### Problem Scenario 2: CNN Architecture Variations

If asked to modify the CNN, remember these patterns:

```
# Adding more conv layers - maintain the pattern
class CustomCNNDeep(nn.Module):
    def __init__(self):
        super().__init__()
        # Pattern: start with 3 channels, double each time
        self.conv1 = nn.Conv2d(3, 32, 3, padding=1)
        self.conv2 = nn.Conv2d(32, 32, 3, padding=1) # Same channels
        self.conv3 = nn.Conv2d(32, 64, 3, padding=1) # Double
        self.conv4 = nn.Conv2d(64, 64, 3, padding=1) # Same
        self.conv5 = nn.Conv2d(64, 128, 3, padding=1) # Double

    def forward(self, x):
        # Pattern: Conv-Conv-Pool, Conv-Conv-Pool
        x = F.relu(self.conv1(x))
        x = F.relu(self.conv2(x))
        x = self.pool(x) # After every 2 conv layers

        x = F.relu(self.conv3(x))
        x = F.relu(self.conv4(x))
        x = self.pool(x)

        # Calculate new flatten size: 128 * 8 * 8 = ____
```

## Scenario 3: Validation vs Training Mode

## Problem Scenario 3: Validation vs Training Mode

Critical differences between training and validation:

```
# Training mode
def train_epoch():
    model.____() # Enable dropout and batch norm training mode

    for batch in train_loader:
        optimizer.____() # Clear gradients
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.____() # Compute gradients
        optimizer.____() # Update weights

# Validation mode
def validate():
    model.____() # Disable dropout, batch norm in eval mode

    with torch.____(): # Disable gradient computation
        for batch in val_loader:
            # NO optimizer.zero_grad() here!
            # NO loss.backward() here!
            # NO optimizer.step() here!
            outputs = model(images)
            loss = criterion(outputs, labels)
```

Scenario 4: RNN with Different Sequence Lengths

## Problem Scenario 4: RNN with Different Sequence Lengths

If given a different text, adapt the RNN code:

```
# Original: "Deep Learning"
text = "Deep Learning"
input_seq = text[:-1] # "Deep Learnin"
target_seq = text[1:] # "eep Learning"

# New text: "Hello World"
text = "Hello World"
input_seq = text[____] # "Hello Worl"
target_seq = text[____] # "ello World"

# Vocabulary size changes!
chars = sorted(list(set(text)))
V = len(chars) # This will be different!

# All weight matrices need to be reinitialized with new V
W_xh = torch.randn(H, ____, requires_grad=True) * 0.1
W_hy = torch.randn(____, H, requires_grad=True) * 0.1
b_y = torch.zeros(____, requires_grad=True)
```

Scenario 5: Hyperparameter Grid Search Pattern

## Problem Scenario 5: Hyperparameter Grid Search Pattern

Standard grid search implementation:



```

learning_rates = [0.0001, 0.001]
optimizers = [torch.optim.Adam, torch.optim.SGD]
model_classes = [CustomCNN, CustomCNNwithBN]

best_accuracy = 0
best_params = None

for model_class in model_classes:
    for optimizer_class in optimizers:
        for lr in learning_rates:
            # CRITICAL: Reinitialize model each time
            model = model_class().to(device)

            # Initialize optimizer based on type
            if optimizer_class == torch.optim.SGD:
                optimizer = optimizer_class(
                    model.parameters(),
                    lr=lr,
                    momentum=____,
                    weight_decay=____
                )
            else: # Adam
                optimizer = optimizer_class(
                    model.parameters(),
                    lr=lr,
                    weight_decay=____
                )

            # Train and validate...
            val_acc = train_and_validate(model, optimizer)

            if val_acc > best_accuracy:
                best_accuracy = val_acc
                best_params = (model_class.__name__, optimizer_class.__name__, lr)

```

Scenario 6: Top-K Accuracy Calculation Variations

## Problem Scenario 6: Top-K Accuracy Calculation Variations

Different ways to calculate accuracy:

```

# Top-1 accuracy (most common)
_, predicted = torch.max(outputs, 1)
correct = (predicted == labels).sum().item()
accuracy = correct / labels.size(0) * 100

# Top-K accuracy using topk
_, top_k_pred = outputs.topk(____, dim=1, largest=True, sorted=True)
top_k_correct = top_k_pred.eq(labels.view(-1, 1).expand_as(____)).sum().item()

# Alternative top-K calculation
values, indices = torch.topk(outputs, k=5, dim=1)
correct_mask = indices == labels.unsqueeze(1)
top_k_accuracy = correct_mask.sum().float() / labels.size(0) * 100

```

# EXAM SIMULATION QUESTIONS

Quick Code Writing

## Problem Quick Code Writing

Write these functions from memory in 2 minutes each:

1. Write the bilinear interpolation weight calculation:

```
def calculate_bilinear_weight(p_x, p_y, q_x, q_y):  
    # Your code here - calculate G(p,q)  
    return ____
```

2. Write the RNN hidden state update:

```
def rnn_step(x_t, h_prev, W_xh, W_hh, b_xh, b_hh):  
    # Your code here - compute new hidden state  
    return ____
```

3. Write the CNN forward pass for one block:

```
def cnn_block_forward(x, conv1, conv2, pool):  
    # Your code here - conv-conv-pool pattern  
    return ____
```

4. Write the training step:

```
def training_step(model, optimizer, criterion, images, labels):  
    # Your code here - complete training step  
    return loss
```

5. Write character to one-hot conversion:

```
def char_to_onehot(char, char2idx, vocab_size):  
    # Your code here - convert character to one-hot vector  
    return ____
```

## FINAL MEMORY CHECK

Last Minute Review

## Problem Last Minute Review

Before the exam, quickly verify you remember:

Constants (write from memory):

- CIFAR-100 mean: [\_\_\_\_, \_\_\_\_, \_\_\_\_]
- CIFAR-100 std: [\_\_\_\_, \_\_\_\_, \_\_\_\_]

- Train/val split ratio: \_\_\_\_
- Dropout probability: \_\_\_\_
- Weight decay: \_\_\_\_
- SGD momentum: \_\_\_\_

**Dimensions (write from memory):**

- CNN channels:  $3 \rightarrow \_\_\_ \rightarrow \_\_\_ \rightarrow \_\_\_ \rightarrow \_\_\_$
- CNN spatial:  $32 \rightarrow \_\_\_ \rightarrow \_\_\_ \rightarrow \_\_\_$
- RNN  $W_{xh}$ : (\_\_\_\_, \_\_\_\_)
- RNN  $W_{hh}$ : (\_\_\_\_, \_\_\_\_)
- RNN  $W_{hy}$ : (\_\_\_\_, \_\_\_\_)

# PRACTICE EXAM - NO PEEKING AT ANSWERS!

## Practice 1: Bilinear Interpolation Implementation

### Problem Practice 1: Bilinear Interpolation Implementation

Complete this bilinear interpolation function:

```
def bilinear_interpolate(a_l, q_y, q_x):
    H, W = a_l.shape

    y0 = int(np._____(q_y))
    x0 = int(np._____(q_x))
    y1 = y0 + ___
    x1 = x0 + ___

    def get_pixel_value(img, y, x):
        if 0 <= y < H and 0 <= x < W:
            return img[y, x]
        else:
            return ____

    v_00 = get_pixel_value(a_l, y0, x0)
    v_01 = get_pixel_value(a_l, y0, x1)
    v_10 = get_pixel_value(a_l, y1, x0)
    v_11 = get_pixel_value(a_l, y1, x1)

    dy = q_y - ___
    dx = q_x - ___

    v_0 = v_00 * (1 - ___) + v_01 * ___
    v_1 = v_10 * (1 - ___) + v_11 * ___
    out = v_0 * (1 - ___) + v_1 * ___

    return out
```

## Practice 2: Deformable Conv Offset Extraction

### Problem Practice 2: Deformable Conv Offset Extraction

Fill in the correct offset extraction pattern:

```
for kh in range(K_h):
    for kw in range(K_w):
        k = kh * K_w + kw

        delta_y = delta[n, ___ * k, h_out, w_out]
        delta_x = delta[n, ___ * k + ___, h_out, w_out]
        m_k = mask[n, k, h_out, w_out]

        sample_y = h_start + kh * dilation + ____
        sample_x = w_start + kw * dilation + ____
```

```
# Apply weight and mask
value += weight[c_out, c_in, kh, kw] * ___ * ____
```

### Practice 3: CNN Architecture

## Problem Practice 3: CNN Architecture

Complete the CustomCNN constructor with correct channel progression:

```
class CustomCNN(nn.Module):
    def __init__(self):
        super(CustomCNN, self).__init__()

        self.conv1 = nn.Conv2d(____, ____, kernel_size=3, padding=1)
        self.conv2 = nn.Conv2d(____, ____, kernel_size=3, padding=1)
        self.conv3 = nn.Conv2d(____, ____, kernel_size=3, padding=1)
        self.conv4 = nn.Conv2d(____, ____, kernel_size=3, padding=1)

        self.pool = nn._____(2, 2)

        self.fc1 = nn.Linear(___ * ___ * ____, 512)
        self.fc2 = nn.Linear(____, ____)
        self.fc3 = nn.Linear(____, ____)

        self.dropout = nn.Dropout(____)
```

### Practice 4: CNN Forward Pass

## Problem Practice 4: CNN Forward Pass

Complete the forward pass with correct activations:

```
def forward(self, x):
    x = F.____(self.conv1(x))
    x = F.____(self.conv2(x))
    x = self.pool(x) # 32x32 -> ____

    x = F.____(self.conv3(x))
    x = self.pool(x) # 16x16 -> ____

    x = F.____(self.conv4(x))
    x = self.pool(x) # 8x8 -> ____

    x = x.view(x.____(), ____)

    x = F.____(self.fc1(x))
    x = self._____(x)
```

```

x = F.____(self.fc2(x))
x = self._____(x)
x = self.fc3(x)

return x

```

Practice 5: Training Loop Pattern

## Problem Practice 5: Training Loop Pattern

Complete the sacred training loop:

```

def train_step(model, optimizer, criterion, images, labels, device):
    images = images.to(____)
    labels = labels.to(____)

    optimizer.____()
    outputs = model(_____)
    loss = criterion(outputs, labels)
    loss._____(____)
    optimizer.____()

    return loss.item()

```

Practice 6: Top-K Accuracy

## Problem Practice 6: Top-K Accuracy

Complete the top-1 and top-5 accuracy calculation:

```

# Top-1 accuracy
_, top1_pred = outputs.topk(____, dim=1, largest=____, sorted=____)
top1_correct = top1_pred.eq(labels.view(____, ____)).sum().item()

# Top-5 accuracy
_, top5_pred = outputs.topk(____, dim=1, largest=____, sorted=____)
top5_correct = top5_pred.eq(labels.view(-1, 1).____(____)).sum().item()

```

Practice 7: RNN Weight Initialization

## Problem Practice 7: RNN Weight Initialization

Initialize RNN parameters with correct dimensions:

```

V = len(chars) # Vocabulary size
H = 16         # Hidden size

W_xh = torch.randn(____, ____, requires_grad=True) * 0.1
W_hh = torch.randn(____, ____, requires_grad=True) * 0.1
b_xh = torch.zeros(____, requires_grad=True)
b_hh = torch.zeros(____, requires_grad=True)
W_hy = torch.randn(____, ____, requires_grad=True) * 0.1
b_y = torch.zeros(____, requires_grad=True)

```

Practice 8: RNN Forward Equations

## Problem Practice 8: RNN Forward Equations

Complete the RNN forward pass:

```

for t in range(seq_len):
    x_t = inputs[t]

    # Hidden state update
    h = torch.tanh(_____ @ x_t + _____ + _____ @ h + _____)

    # Output logits
    s_t = _____ @ h + _____

    logits_list.append(s_t)

```

Practice 9: Character Processing

## Problem Practice 9: Character Processing

Complete the character-level preprocessing:

```

text = "Deep Learning"

chars = sorted(list(set(_____)))
char2idx = {ch: i for i, ch in enumerate(_____)}
idx2char = {i: ch for i, ch in enumerate(_____)}

input_seq = text[_____]
target_seq = text[_____]

def one_hot(idx, size):
    vec = torch.zeros(size)
    vec[_____] = 1.0
    return vec

```

```
inputs = [one_hot(char2idx[ch], V) for ch in ____]
targets = torch.tensor([char2idx[ch] for ch in ____], dtype=torch.long)
```

## Practice 10: Gradient Computation

### Problem Practice 10: Gradient Computation

Complete the explicit gradient computation:

```
grad_W_xh = torch.autograd.grad(loss_manual, W_xh, retain_graph=____)[0]
grad_W_hh = torch.autograd.grad(loss_manual, W_hh, retain_graph=____)[0]
grad_b_xh = torch.autograd.grad(loss_manual, b_xh, retain_graph=____)[0]
grad_b_hh = torch.autograd.grad(loss_manual, b_hh, retain_graph=____)[0]
grad_W_hy = torch.autograd.grad(loss_manual, W_hy, retain_graph=____)[0]
grad_b_y = torch.autograd.grad(loss_manual, b_y, retain_graph=____)[0]
```

## Practice 11: BatchNorm Integration

### Problem Practice 11: BatchNorm Integration

Add BatchNorm to CNN architecture:

```
class CustomCNNwithBN(nn.Module):
    def __init__(self):
        super().__init__()

        self.conv1 = nn.Conv2d(3, 32, kernel_size=3, padding=1)
        self.bn1 = nn.BatchNorm2d(____)

        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, padding=1)
        self.bn2 = nn.BatchNorm2d(____)

    def forward(self, x):
        x = F.relu(self.____(self.conv1(x)))
        x = F.relu(self.____(self.conv2(x)))
```

## Practice 12: CIFAR-100 Constants

### Problem Practice 12: CIFAR-100 Constants

Fill in the CIFAR-100 normalization values:

```
transform = transforms.Compose([
```



```

        transforms.ToTensor(),
        transforms.Normalize(mean=[____, ____, ____],
                               std=[____, ____, ____])
    ])

```

### Practice 13: Data Splitting

## Problem Practice 13: Data Splitting

Complete the 80/20 train/validation split:

```

full_train_set = CIFAR100(root='./data', train=True, download=True)

train_size = int(____ * len(full_train_set))
val_size = len(full_train_set) - ____

train_set, val_set = random_split(full_train_set, [____, ____])

```

### Practice 14: Loss and Optimizer Setup

## Problem Practice 14: Loss and Optimizer Setup

Set up loss function and optimizer for CIFAR-100:

```

loss_function = nn.______()

optimizer = optim.SGD(model.parameters(),
                       lr=____,
                       momentum=____,
                       weight_decay=____)

```

### Practice 15: Quick Recall

## Problem Practice 15: Quick Recall

Answer these without looking:

1. What function rounds down to nearest integer? \_\_\_\_\_
2. In PyTorch deformable conv, Y offset comes at index \_\_\_\_ and X offset at \_\_\_\_
3. The bilinear interpolation weight formula is:  $G(p,q) = \text{_____}$

4. CNN channel progression in our model:  $3 \rightarrow \_\_\_ \rightarrow \_\_\_ \rightarrow \_\_\_ \rightarrow \_\_\_$
5. RNN hidden state uses which activation function? \_\_\_\_\_
6. In training loop, what comes after `loss.backward()`? \_\_\_\_\_
7. For validation, you wrap the loop with `torch._____()`:
8. The flatten operation uses `x._____(x._____(0), ____)`:
9. Top-5 accuracy uses `outputs._____(____, dim=1)`:
10. RNN weight `W_hy` has dimensions `(____, ____)`:

## CHALLENGE PROBLEMS

Challenge 1: Complete Function Implementation

### **Problem Challenge 1: Complete Function Implementation**

Write the complete bilinear interpolation function from scratch:

Challenge 2: Complete RNN Step

### **Problem Challenge 2: Complete RNN Step**

Write a complete RNN forward step function:

Challenge 3: Training vs Validation

### **Problem Challenge 3: Training vs Validation**

Write the key differences between training and validation loops:

#### Challenge 4: Debugging Scenario

### Problem Challenge 4: Debugging Scenario

Your deformable convolution outputs don't match PyTorch. List 5 most likely bugs:

1. \_\_\_\_\_
2. \_\_\_\_\_
3. \_\_\_\_\_
4. \_\_\_\_\_
5. \_\_\_\_\_

#### Challenge 5: Architecture Design

### Problem Challenge 5: Architecture Design

Design a deeper CNN with 6 conv layers following the same pattern:

## ANSWER BANK

Practice 1 Answers:

### Problem Practice 1 Answers:

floor, floor, 1, 1, 0.0, y0, x0, dx, dx, dx, dx, dy, dy

Practice 2 Answers:

### Problem Practice 2 Answers:

2, 2, 1, delta\_y, delta\_x, m\_k, interpolated

Practice 3 Answers:

### Problem Practice 3 Answers:

3, 32, 32, 64, 64, 128, 128, 256, MaxPool2d, 256, 4, 4, 512, 256, 256, 100, 0.5

Practice 4 Answers:

### Problem Practice 4 Answers:

relu, relu, 16x16, relu, 8x8, relu, 4x4, size, 0, -1, relu, dropout, relu, dropout

Practice 5 Answers:

### Problem Practice 5 Answers:

device, device, zero\_grad, images, backward, step

Practice 6 Answers:

### Problem Practice 6 Answers:

1, True, True, -1, 1, 5, True, True, expand\_as(top5\_pred)

Practice 7 Answers:

### Problem Practice 7 Answers:

H, V, H, H, H, H, V, H, V

Practice 8 Answers:

### Problem Practice 8 Answers:

W\_xh, b\_xh, W\_hh, b\_hh, W\_hy, b\_y

Practice 9 Answers:

### Problem Practice 9 Answers:

text, chars, chars, [:-1], [1:], idx, input\_seq, target\_seq

Practice 10 Answers:

### Problem Practice 10 Answers:

True, True, True, True, True, True

Practice 11 Answers:

### Problem Practice 11 Answers:

32, 64, bn1, bn2

Practice 12 Answers:

### Problem Practice 12 Answers:

0.5071, 0.4867, 0.4408, 0.2675, 0.2565, 0.2761

Practice 13 Answers:

### Problem Practice 13 Answers:

0.8, train\_size, train\_size, val\_size

Practice 14 Answers:

### Problem Practice 14 Answers:

CrossEntropyLoss, 0.01, 0.9, 5e-4

Practice 15 Answers:

### Problem Practice 15 Answers:

1. np.floor 2. 2\*k, 2\*k+1 3.  $(1 - p_x - q_x) * (1 - p_y - q_y)$  4. 32, 64, 128, 256 5. tanh 6. optimizer.step()  
7. no\_grad 8. view, size, -1 9. topk, 5 10. V, H

Challenge Problem Answers:

### Problem Challenge Problem Answers:

**Challenge 1:** Complete bilinear function with all bounds checking, corner calculations, and interpolation steps.

**Challenge 2:** RNN step with tanh activation: `h = torch.tanh(W_xh @ x + b_xh + W_hh @ h_prev + b_hh)`

**Challenge 3:** Training: `model.train()`, optimizer steps, gradients enabled. Validation: `model.eval()`, `torch.no_grad()`, no optimizer steps.

**Challenge 4:** 1. Wrong offset indexing (2\*k), 2. Wrong bilinear order (X first), 3. Missing zero padding, 4. Wrong sampling position calculation, 5. Incorrect weight/mask application.

**Challenge 5:** Conv layers: 3→32→32→64→64→128→128, with pooling after every 2 conv layers.

# CENG403 - Spring 2025: Homework set Write Code Blocks Practice

Your Name

## TASK 1: DEFORMABLE CNN - WRITE CODE BLOCKS

.1: Floor Operation

### Problem 1.1: Floor Operation

Write code to get the floor of a float value  $q_y$  and store it in  $y0$ :

.2: Ceiling Operation

### Problem 1.2: Ceiling Operation

Write code to get the ceiling of a float value  $q_x$  and store it in  $x1$ :

.3: Kernel Index Calculation

### Problem 1.3: Kernel Index Calculation

Write code to calculate linear kernel index from 2D position  $(kh, kw)$  with kernel width  $K_w$ :

.4: Y Offset Extraction

### Problem 1.4: Y Offset Extraction

Write code to extract Y offset from delta tensor for batch  $n$ , kernel index  $k$ , output position  $(h_{out}, w_{out})$ :

.5: X Offset Extraction

### Problem 1.5: X Offset Extraction

Write code to extract X offset from delta tensor for batch n, kernel index k, output position (h\_out, w\_out):

.6: Base Position Calculation

### Problem 1.6: Base Position Calculation

Write code to calculate base sampling position h\_start from output position h\_out and stride:

.7: Final Sampling Position

### Problem 1.7: Final Sampling Position

Write code to calculate final sampling position sample\_y from h\_start, kh, dilation, and delta\_y:

.8: Bounds Check Condition

### Problem 1.8: Bounds Check Condition

Write code to check if coordinates (y, x) are within image bounds (H, W):

.9: Safe Pixel Access

### Problem 1.9: Safe Pixel Access

Write code to get pixel value at (y, x) from image img, returning 0.0 if out of bounds:



.10: Fractional Part Calculation

### **Problem 1.10: Fractional Part Calculation**

Write code to calculate fractional part  $dx$  from  $q_x$  and its floor  $x0$ :

.11: Linear Interpolation

### **Problem 1.11: Linear Interpolation**

Write code to linearly interpolate between  $v_{left}$  and  $v_{right}$  using weight  $dx$ :

.12: Bilinear Weight Calculation

### **Problem 1.12: Bilinear Weight Calculation**

Write code to calculate bilinear interpolation weight for points  $p$  and  $q$ :

.13: Value Accumulation

### **Problem 1.13: Value Accumulation**

Write code to accumulate weighted interpolated value to  $current\_value$  using  $weight$ ,  $mask$ , and interpolated:

.14: Corner Position Calculation

### **Problem 1.14: Corner Position Calculation**

Write code to calculate all four corner positions  $(y_0, x_0)$ ,  $(y_0, x_1)$ ,  $(y_1, x_0)$ ,  $(y_1, x_1)$  from  $q_y$  and  $q_x$ :

.15: Mask Application

### **Problem 1.15: Mask Application**

Write code to apply modulation mask  $m_k$  to an interpolated value:

## TASK 2: CNN PYTORCH - WRITE CODE BLOCKS

.1: Conv Layer Definition

### Problem 2.1: Conv Layer Definition

Write code to define a Conv2d layer with 32 input channels, 64 output channels, kernel size 3, padding 1:

.2: BatchNorm Layer Definition

### Problem 2.2: BatchNorm Layer Definition

Write code to define a BatchNorm2d layer for 128 channels:

.3: MaxPool Layer Definition

### Problem 2.3: MaxPool Layer Definition

Write code to define a MaxPool2d layer that halves spatial dimensions:

.4: Linear Layer Definition

### Problem 2.4: Linear Layer Definition

Write code to define a Linear layer with 4096 inputs and 512 outputs:

.5: Dropout Layer Definition

### Problem 2.5: Dropout Layer Definition

Write code to define a Dropout layer with probability 0.5:

.6: ReLU Activation

### **Problem 2.6: ReLU Activation**

Write code to apply ReLU activation to variable x:

.7: Tensor Flattening

### **Problem 2.7: Tensor Flattening**

Write code to flatten tensor x while preserving batch dimension:

.8: Device Transfer

### **Problem 2.8: Device Transfer**

Write code to move tensor images to device:

.9: Optimizer Zero Grad

### **Problem 2.9: Optimizer Zero Grad**

Write code to clear gradients from optimizer:

.10: Forward Pass

### **Problem 2.10: Forward Pass**

Write code to perform forward pass through model with input images:

.11: Loss Calculation

### **Problem 2.11: Loss Calculation**

Write code to calculate loss using criterion with outputs and labels:

.12: Backward Pass

### **Problem 2.12: Backward Pass**

Write code to perform backward pass on loss:

.13: Optimizer Step

### **Problem 2.13: Optimizer Step**

Write code to update model parameters using optimizer:

.14: Top-1 Prediction

### **Problem 2.14: Top-1 Prediction**

Write code to get top-1 predictions from outputs:

.15: Accuracy Calculation

### **Problem 2.15: Accuracy Calculation**

Write code to calculate accuracy from predicted and labels tensors:

.16: Model Training Mode

### **Problem 2.16: Model Training Mode**

Write code to set model to training mode:

.17: Model Evaluation Mode

### **Problem 2.17: Model Evaluation Mode**

Write code to set model to evaluation mode:

.18: No Gradient Context

### **Problem 2.18: No Gradient Context**

Write code to create context where gradients are disabled:

.19: Dataset Size Calculation

### **Problem 2.19: Dataset Size Calculation**

Write code to calculate training size as 80% of total dataset size:

.20: Random Split

### **Problem 2.20: Random Split**

Write code to split dataset into `train_set` and `val_set` with sizes `train_size` and `val_size`:

## TASK 3: RNN - WRITE CODE BLOCKS

.1: Character Set Creation

### Problem 3.1: Character Set Creation

Write code to create sorted list of unique characters from text:

.2: Character to Index Mapping

### Problem 3.2: Character to Index Mapping

Write code to create dictionary mapping each character to its index:

.3: Index to Character Mapping

### Problem 3.3: Index to Character Mapping

Write code to create dictionary mapping each index to its character:

.4: Input Sequence Creation

### Problem 3.4: Input Sequence Creation

Write code to create input sequence (all characters except last) from text:



.5: Target Sequence Creation

### **Problem 3.5: Target Sequence Creation**

Write code to create target sequence (all characters except first) from text:

.6: One-Hot Vector Creation

### **Problem 3.6: One-Hot Vector Creation**

Write code to create one-hot vector of given size with 1 at given index:

.7: Input-to-Hidden Weight Initialization

### **Problem 3.7: Input-to-Hidden Weight Initialization**

Write code to initialize  $W_{xh}$  weight matrix for RNN with hidden size  $H$  and vocab size  $V$ :

.8: Hidden-to-Hidden Weight Initialization

### **Problem 3.8: Hidden-to-Hidden Weight Initialization**

Write code to initialize  $W_{hh}$  weight matrix for RNN with hidden size  $H$ :

.9: Hidden Bias Initialization

### **Problem 3.9: Hidden Bias Initialization**

Write code to initialize bias vector  $b_{xh}$  for hidden layer with size  $H$ :

.10: Output Weight Initialization

### **Problem 3.10: Output Weight Initialization**

Write code to initialize  $W_{hy}$  weight matrix from hidden to output with vocab size  $V$  and hidden size  $H$ :

.11: Hidden State Initialization

### **Problem 3.11: Hidden State Initialization**

Write code to initialize hidden state vector with zeros of size  $H$ :

.12: Input Contribution Calculation

### **Problem 3.12: Input Contribution Calculation**

Write code to calculate input contribution to hidden state using  $W_{xh}$  and  $x_t$ :

.13: Hidden Recurrence Calculation

### **Problem 3.13: Hidden Recurrence Calculation**

Write code to calculate recurrent contribution to hidden state using  $W_{hh}$  and previous hidden state  $h$ :

.14: Hidden State Update

### Problem 3.14: Hidden State Update

Write code to update hidden state using  $\tanh$  activation with all contributions and biases:

.15: Output Logits Calculation

### Problem 3.15: Output Logits Calculation

Write code to calculate output logits  $s_t$  from hidden state  $h$  using  $W_{hy}$  and  $b_y$ :

.16: Character Index Lookup

### Problem 3.16: Character Index Lookup

Write code to get index of character 'e' from `char2idx` dictionary:

.17: Input List Creation

### Problem 3.17: Input List Creation

Write code to convert input sequence to list of one-hot vectors:

.18: Target Tensor Creation

### **Problem 3.18: Target Tensor Creation**

Write code to convert target sequence to tensor of character indices:

.19: Logits Stacking

### **Problem 3.19: Logits Stacking**

Write code to stack list of logits tensors into single tensor:

.20: Log Softmax Calculation

### **Problem 3.20: Log Softmax Calculation**

Write code to calculate log softmax of logits along vocabulary dimension:

.21: NLL Loss Calculation

### **Problem 3.21: NLL Loss Calculation**

Write code to calculate negative log likelihood loss from log\_probs and targets:

.22: Gradient Calculation

### **Problem 3.22: Gradient Calculation**

Write code to calculate gradient of loss with respect to  $W_{xh}$ :

## DATA PREPROCESSING - WRITE CODE BLOCKS

.1: CIFAR-100 Mean Values

### Problem 4.1: CIFAR-100 Mean Values

Write code to define CIFAR-100 normalization mean values:

.2: CIFAR-100 Std Values

### Problem 4.2: CIFAR-100 Std Values

Write code to define CIFAR-100 normalization standard deviation values:

.3: ToTensor Transform

### Problem 4.3: ToTensor Transform

Write code to create ToTensor transform:

.4: Normalization Transform

### Problem 4.4: Normalization Transform

Write code to create Normalize transform with CIFAR-100 mean and std:

.5: Random Crop Transform

## Problem 4.5: Random Crop Transform

Write code to create RandomCrop transform with size 32 and padding 4:

.6: Random Flip Transform

## Problem 4.6: Random Flip Transform

Write code to create RandomHorizontalFlip transform:

.7: Transform Composition

## Problem 4.7: Transform Composition

Write code to compose multiple transforms into single transform:

.8: CIFAR-100 Dataset Loading

## Problem 4.8: CIFAR-100 Dataset Loading

Write code to load CIFAR-100 training dataset with transform:

.9: DataLoader Creation

## Problem 4.9: DataLoader Creation

Write code to create DataLoader with batch size 128 and shuffle=True:



## OPTIMIZATION AND TRAINING - WRITE CODE BLOCKS

.1: CrossEntropy Loss Definition

### Problem 5.1: CrossEntropy Loss Definition

Write code to define CrossEntropy loss function:

.2: SGD Optimizer Definition

### Problem 5.2: SGD Optimizer Definition

Write code to define SGD optimizer with learning rate 0.01 and momentum 0.9:

.3: Adam Optimizer Definition

### Problem 5.3: Adam Optimizer Definition

Write code to define Adam optimizer with learning rate 0.001:

.4: Model Parameter Count

### Problem 5.4: Model Parameter Count

Write code to count total number of parameters in model:

.5: Learning Rate Update

### Problem 5.5: Learning Rate Update

Write code to multiply learning rate by 0.1 for all parameter groups:

.6: Model State Save

### Problem 5.6: Model State Save

Write code to save model state dictionary to file 'model.pth':

.7: Model State Load

### Problem 5.7: Model State Load

Write code to load model state dictionary from file 'model.pth':

.8: Gradient Clipping

### Problem 5.8: Gradient Clipping

Write code to clip gradients to maximum norm of 1.0:

3cm

.9: Top-5 Accuracy

### Problem 5.9: Top-5 Accuracy

Write code to calculate top-5 accuracy from outputs and labels:

.10: Loss Item Extraction

## Problem 5.10: Loss Item Extraction

Write code to extract scalar value from loss tensor:

## DEBUGGING AND UTILITIES - WRITE CODE BLOCKS

.1: Tensor Shape Check

### Problem 6.1: Tensor Shape Check

Write code to print shape of tensor x:

.2: Tensor Device Check

### Problem 6.2: Tensor Device Check

Write code to check which device tensor x is on:

.3: Model Device Transfer

### Problem 6.3: Model Device Transfer

Write code to move entire model to GPU:

.4: Gradient Existence Check

### Problem 6.4: Gradient Existence Check

Write code to check if parameter has gradients:

.5: Memory Usage Check

### Problem 6.5: Memory Usage Check

Write code to check CUDA memory usage:

.6: Random Seed Setting

### **Problem 6.6: Random Seed Setting**

Write code to set PyTorch random seed to 42:

.7: Numpy Seed Setting

### **Problem 6.7: Numpy Seed Setting**

Write code to set numpy random seed to 42:

.8: Model Summary

### **Problem 6.8: Model Summary**

Write code to print model architecture:

.9: Batch Dimension Check

### **Problem 6.9: Batch Dimension Check**

Write code to get batch size from tensor x:

.10: Tensor Type Conversion

### **Problem 6.10: Tensor Type Conversion**

Write code to convert tensor x to float type:

## ANSWER BANK

Task 1 - Deformable CNN Answers:

### Problem Task 1 - Deformable CNN Answers:

```
1.1: y0 = int(np.floor(q_y))
1.2: x1 = int(np.ceil(q_x))
1.3: k = kh * K_w + kw
1.4: delta_y = delta[n, 2 * k, h_out, w_out]
1.5: delta_x = delta[n, 2 * k + 1, h_out, w_out]
1.6: h_start = h_out * stride
1.7: sample_y = h_start + kh * dilation + delta_y
1.8: if 0 <= y < H and 0 <= x < W:
1.9: value = img[y, x] if (0 <= y < H and 0 <= x < W) else 0.0
1.10: dx = q_x - x0
1.11: result = v_left * (1 - dx) + v_right * dx
1.12: weight = (1 - abs(p_x - q_x)) * (1 - abs(p_y - q_y))
1.13: current_value += weight * mask * interpolated
1.14: y0, x0 = int(np.floor(q_y)), int(np.floor(q_x))
    y1, x1 = y0 + 1, x0 + 1
1.15: modulated_value = m_k * interpolated_value
```

Task 2 - CNN PyTorch Answers:

### Problem Task 2 - CNN PyTorch Answers:

```
2.1: self.conv = nn.Conv2d(32, 64, kernel_size=3, padding=1)
2.2: self.bn = nn.BatchNorm2d(128)
2.3: self.pool = nn.MaxPool2d(2, 2)
2.4: self.fc = nn.Linear(4096, 512)
2.5: self.dropout = nn.Dropout(0.5)
2.6: x = F.relu(x)
2.7: x = x.view(x.size(0), -1)
2.8: images = images.to(device)
2.9: optimizer.zero_grad()
2.10: outputs = model(images)
2.11: loss = criterion(outputs, labels)
```

```

2.12: loss.backward()

2.13: optimizer.step()

2.14: _, predicted = torch.max(outputs, 1)

2.15: accuracy = (predicted == labels).sum().item() / labels.size(0) * 100

2.16: model.train()

2.17: model.eval()

2.18: with torch.no_grad():

2.19: train_size = int(0.8 * len(dataset))

2.20: train_set, val_set = random_split(dataset, [train_size, val_size])

```

Task 3 - RNN Answers:

## Problem Task 3 - RNN Answers:

```

3.1: chars = sorted(list(set(text)))

3.2: char2idx = {ch: i for i, ch in enumerate(chars)}

3.3: idx2char = {i: ch for i, ch in enumerate(chars)}

3.4: input_seq = text[:-1]

3.5: target_seq = text[1:]

3.6: vec = torch.zeros(size)
vec[idx] = 1.0

3.7: W_xh = torch.randn(H, V, requires_grad=True) * 0.1

3.8: W_hh = torch.randn(H, H, requires_grad=True) * 0.1

3.9: b_xh = torch.zeros(H, requires_grad=True)

3.10: W_hy = torch.randn(V, H, requires_grad=True) * 0.1

3.11: h = torch.zeros(H)

3.12: input_contrib = W_xh @ x_t

3.13: hidden_contrib = W_hh @ h

3.14: h = torch.tanh(W_xh @ x_t + b_xh + W_hh @ h + b_hh)

3.15: s_t = W_hy @ h + b_y

3.16: idx = char2idx['e']

3.17: inputs = [one_hot(char2idx[ch], V) for ch in input_seq]

3.18: targets = torch.tensor([char2idx[ch] for ch in target_seq], dtype=torch.long)

3.19: logits = torch.stack(logits_list)

3.20: log_probs = F.log_softmax(logits, dim=1)

```



```
3.21: loss = F.nll_loss(log_probs, targets)
```

```
3.22: grad.W_xh = torch.autograd.grad(loss, W_xh, retain_graph=True)[0]
```

Additional Sections Answers:

## Problem Additional Sections Answers:

```
4.1: mean = [0.5071, 0.4867, 0.4408]
```

```
4.2: std = [0.2675, 0.2565, 0.2761]
```

```
4.3: transform = transforms.ToTensor()
```

```
4.4: normalize = transforms.Normalize(mean=[0.5071, 0.4867, 0.4408], std=[0.2675, 0.2565, 0.2761])
```

```
4.5: crop = transforms.RandomCrop(32, padding=4)
```

```
4.6: flip = transforms.RandomHorizontalFlip()
```

```
4.7: transform = transforms.Compose([transform1, transform2, ...])
```

```
4.8: dataset = CIFAR100(root='./data', train=True, transform=transform)
```

```
4.9: loader = DataLoader(dataset, batch_size=128, shuffle=True)
```

```
5.1: criterion = nn.CrossEntropyLoss()
```

```
5.2: optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9)
```

```
5.3: optimizer = optim.Adam(model.parameters(), lr=0.001)
```

```
5.4: total_params = sum(p.numel() for p in model.parameters())
```

```
5.5: for param_group in optimizer.param_groups: param_group['lr'] *= 0.1
```

```
5.6: torch.save(model.state_dict(), 'model.pth')
```

```
5.7: model.load_state_dict(torch.load('model.pth'))
```

```
5.8: torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)
```

```
5.9: _, top5_pred = outputs.topk(5, dim=1)
top5_acc = top5_pred.eq(labels.view(-1, 1).expand_as(top5_pred)).sum().item()
```

```
5.10: loss_value = loss.item()
```

```
6.1: print(x.shape)
```

```
6.2: print(x.device)
```

```
6.3: model = model.to('cuda')
```

```
6.4: if param.grad is not None:
```

```
6.5: print(torch.cuda.memory_allocated())
```

```
6.6: torch.manual_seed(42)
```

```
6.7: np.random.seed(42)
```

```
6.8: print(model)
```

6.9: `batch_size = x.size(0)`

6.10: `x = x.float()`

Additional Answers

## Problem Additional Answers

C1: 32 x 32, C2: 8 x 8, C3: 4096, C4: 73856, C5: 256

S1: `numpy`, `torch.nn`, `F`, `transforms`

S2: `nn.Module`, `super`

S3: `def forward(self, x)`

S4: `nn.CrossEntropyLoss()`

S5: `optim.SGD(model.parameters(), lr=0.01, momentum=0.9)`

M1: `x = self.dropout(x)` if `self.training` else `x`

M2: `for param_group in optimizer.param_groups: param_group['lr'] *= 0.1`

M3: `torch.save(model.state_dict(), 'model.pth')`

M4: `outputs = model(images.to(device))`

## ADVANCED DEFORMABLE CNN OPERATIONS

D1

### Problem D1

Write code to calculate output dimensions after deformable convolution:

D2

### Problem D2

Write code to pad input tensor with zeros on all sides by padding amount:

D3

### Problem D3

Write code to initialize output tensor for deformable convolution with correct shape:

D4

### Problem D4

Write code to extract four corner values for bilinear interpolation given positions:

D5

### Problem D5

Write code to apply dilation to kernel position in deformable convolution:

D6

### Problem D6

Write code to accumulate convolution result across all input channels:

D7

### Problem D7

Write code to validate that sampling position is fractional:

D8

### Problem D8

Write code to compute weighted sum for deformable convolution at one position:

## ADVANCED CNN OPERATIONS

A1

### Problem A1

Write code to calculate receptive field size after multiple conv layers:

A2

### Problem A2

Write code to implement skip connection (residual connection):

A3

### Problem A3

Write code to apply different transforms for training vs validation:

A4

### Problem A4

Write code to calculate model memory usage:

A5

### **Problem A5**

Write code to freeze specific layers during training:

A6

### **Problem A6**

Write code to implement learning rate scheduling:

A7

### **Problem A7**

Write code to calculate class-wise accuracy:

A8

### **Problem A8**

Write code to implement early stopping condition:

A9

### Problem A9

Write code to apply different dropout rates for different layers:

A10

### Problem A10

Write code to implement gradient accumulation for large batches:

## ADVANCED RNN OPERATIONS

R1

### Problem R1

Write code to handle variable length sequences in RNN:

R2

### Problem R2

Write code to implement bidirectional RNN forward pass:

R3

### Problem R3

Write code to calculate perplexity from RNN loss:

R4

### Problem R4

Write code to implement teacher forcing during training:



R5

### **Problem R5**

Write code to generate text using trained RNN:

R6

### **Problem R6**

Write code to implement attention mechanism for RNN:

R7

### **Problem R7**

Write code to handle padding in RNN sequences:

R8

### **Problem R8**

Write code to implement LSTM cell from scratch:

R9

## Problem R9

Write code to calculate gradient flow through time steps:

R10

## Problem R10

Write code to implement sequence-to-sequence mapping:

## ERROR HANDLING AND DEBUGGING

E1

### Problem E1

Write code to check if tensors are on the same device:

E2

### Problem E2

Write code to handle CUDA out of memory error:

E3

### Problem E3

Write code to validate input tensor shapes before processing:

E4

### Problem E4

Write code to check for NaN values in gradients:

E5

### Problem E5

Write code to log training progress every N steps:

E6

### Problem E6

Write code to handle empty batch gracefully:

E7

### Problem E7

Write code to validate model output dimensions:

E8

### Problem E8

Write code to catch and handle gradient explosion:

E9

### Problem E9

Write code to verify learning rate is positive:

E10

### Problem E10

Write code to check if model is in correct mode for evaluation:

## MODEL EVALUATION AND METRICS

V1

### Problem V1

Write code to calculate precision for multi-class classification:

V2

### Problem V2

Write code to calculate recall for specific class:

V3

### Problem V3

Write code to compute F1-score from precision and recall:

V4

### Problem V4

Write code to create confusion matrix:

V5

## Problem V5

Write code to calculate mean average precision:

V6

## Problem V6

Write code to implement cross-validation split:

V7

## Problem V7

Write code to calculate model inference time:

V8

## Problem V8

Write code to compute per-class accuracy:

V9

## Problem V9

Write code to calculate balanced accuracy:

V10

## Problem V10

Write code to evaluate model on subset of classes:



## DATA LOADING VARIATIONS

L1

### Problem L1

Write code to create weighted sampler for imbalanced dataset:

L2

### Problem L2

Write code to implement custom collate function:

L3

### Problem L3

Write code to handle corrupted data samples:

L4

### Problem L4

Write code to implement data augmentation pipeline:

L5

### **Problem L5**

Write code to create stratified train/val split:

L6

### **Problem L6**

Write code to implement multi-scale image loading:

L7

### **Problem L7**

Write code to balance dataset using oversampling:

L8

### **Problem L8**

Write code to implement k-fold cross validation data split:

L9

## Problem L9

Write code to create data loader with custom worker init:

L10

## Problem L10

Write code to implement online data augmentation:

# OPTIMIZATION TECHNIQUES

O1

## Problem O1

Write code to implement cosine annealing learning rate:

O2

## Problem O2

Write code to add L1 regularization to loss:

O3

## Problem O3

Write code to implement momentum SGD from scratch:

O4

## Problem O4

Write code to apply different learning rates to different layers:

O5

## Problem O5

Write code to implement AdamW optimizer setup:

O6

## Problem O6

Write code to implement linear warmup schedule:

O7

## Problem O7

Write code to add noise to gradients:

O8

## Problem O8

Write code to implement cyclical learning rates:

O9

## Problem O9

Write code to calculate effective learning rate:

O10

## Problem O10

Write code to implement gradient centralization:

## TENSOR OPERATIONS

T1

### Problem T1

Write code to reshape tensor while preserving total elements:

T2

### Problem T2

Write code to concatenate tensors along specific dimension:

T3

### Problem T3

Write code to split tensor into equal chunks:

T4

### Problem T4

Write code to transpose last two dimensions:

T5

### Problem T5

Write code to compute element-wise maximum of two tensors:

T6

### Problem T6

Write code to select top-k elements along dimension:

T7

### Problem T7

Write code to create mask for padding tokens:

T8

### Problem T8

Write code to compute pairwise distances between vectors:



T9

### Problem T9

Write code to normalize tensor to unit length:

T10

### Problem T10

Write code to apply sliding window operation:

## MEMORY AND PERFORMANCE

P1

### Problem P1

Write code to enable mixed precision training:

P2

### Problem P2

Write code to clear GPU cache:

P3

### Problem P3

Write code to profile memory usage:

P4

### Problem P4

Write code to implement checkpointing for memory efficiency:

P5

### Problem P5

Write code to use `torch.no_grad()` for inference :

P6

### Problem P6

Write code to pin memory for faster data loading:

P7

### Problem P7

Write code to set number of threads for CPU operations:

P8

### Problem P8

Write code to benchmark model inference speed:

P9

### **Problem P9**

Write code to implement gradient checkpointing:

P10

### **Problem P10**

Write code to optimize model for inference:

## EXTENDED ANSWER BANK

Advanced Deformable CNN Answers:

### Problem Advanced Deformable CNN Answers:

```
D1: H_out = (H_in + 2*padding - dilation*(K_h-1) - 1) // stride + 1
D2: padded_input = F.pad(input, (padding, padding, padding, padding))
D3: output = np.zeros((N, C_out, H_out, W_out), dtype=np.float32)
D4: v_00, v_01 = get_pixel(img, y0, x0), get_pixel(img, y0, x1)
    v_10, v_11 = get_pixel(img, y1, x0), get_pixel(img, y1, x1)
D5: dilated_pos_y = h.start + kh * dilation
D6: for c_in in range(C_in): value += weight[c_out, c_in, kh, kw] * interpolated[c_in]
D7: assert sample_y != int(sample_y) or sample_x != int(sample_x)
D8: result = sum(w[k] * m[k] * bilinear_interp(input, pos[k]) for k in range(K))
```

Advanced CNN Answers:

### Problem Advanced CNN Answers:

```
A1: receptive_field = ((kernel_size - 1) * dilation + 1)
A2: x = F.relu(self.conv(x) + x) # residual connection
A3: transform = train_transform if self.training else val_transform
A4: memory_usage = sum(p.numel() * p.element_size() for p in model.parameters())
A5: for param in model.layer.parameters(): param.requires_grad = False
A6: scheduler.step(); current_lr = scheduler.get_last_lr()[0]
A7: per_class_acc = [(pred==i).sum()/(labels==i).sum() for i in range(num_classes)]
A8: if val_loss > best_loss + patience_delta: stop_training = True
A9: self.dropout1 = nn.Dropout(0.2); self.dropout2 = nn.Dropout(0.5)
A10: if (step + 1) % accumulation_steps == 0: optimizer.step(); optimizer.zero_grad()
```

Advanced RNN Answers:

### Problem Advanced RNN Answers:

```
R1: packed_seq = nn.utils.rnn.pack_padded_sequence(x, lengths, batch_first=True)
R2: h_forward = rnn_forward(x); h_backward = rnn_backward(x[:,::-1])
R3: perplexity = torch.exp(loss)
R4: decoder_input = target[:-1] if training else previous_output
R5: with torch.no_grad(): output = model.generate(start_token, max_length)
R6: attention_weights = F.softmax(torch.matmul(query, keys.T), dim=-1)
```

```
R7: mask = (sequence != pad_token).float().unsqueeze(-1)
```

```
R8: f_gate = torch.sigmoid(W_f @ x + U_f @ h + b_f)
```

```
R9: grad_h = torch.autograd.grad(loss, hidden_states, retain_graph=True)
```

```
R10: decoder_output = decoder(encoder_output, target_sequence)
```

Additional Sections Available Upon Request...

**Problem Additional Sections Available Upon Request...**



# CENG403 - Spring 2025: Homework set Write Code Blocks Practice

Your Name

## TASK 1: DEFORMABLE CNN - WRITE CODE BLOCKS

.1: Floor Operation

### Problem 1.1: Floor Operation

Write code to get the floor of a float value  $q_y$  and store it in  $y0$ :

.2: Ceiling Operation

### Problem 1.2: Ceiling Operation

Write code to get the ceiling of a float value  $q_x$  and store it in  $x1$ :

.3: Kernel Index Calculation

### Problem 1.3: Kernel Index Calculation

Write code to calculate linear kernel index from 2D position  $(kh, kw)$  with kernel width  $K_w$ :

.4: Y Offset Extraction

### Problem 1.4: Y Offset Extraction

Write code to extract Y offset from delta tensor for batch  $n$ , kernel index  $k$ , output position  $(h_{out}, w_{out})$ :

.5: X Offset Extraction

### Problem 1.5: X Offset Extraction

Write code to extract X offset from delta tensor for batch n, kernel index k, output position (h\_out, w\_out):

.6: Base Position Calculation

### Problem 1.6: Base Position Calculation

Write code to calculate base sampling position h\_start from output position h\_out and stride:

.7: Final Sampling Position

### Problem 1.7: Final Sampling Position

Write code to calculate final sampling position sample\_y from h\_start, kh, dilation, and delta\_y:

.8: Bounds Check Condition

### Problem 1.8: Bounds Check Condition

Write code to check if coordinates (y, x) are within image bounds (H, W):

.9: Safe Pixel Access

### Problem 1.9: Safe Pixel Access

Write code to get pixel value at (y, x) from image img, returning 0.0 if out of bounds:



.10: Fractional Part Calculation

### **Problem 1.10: Fractional Part Calculation**

Write code to calculate fractional part  $dx$  from  $q_x$  and its floor  $x0$ :

.11: Linear Interpolation

### **Problem 1.11: Linear Interpolation**

Write code to linearly interpolate between  $v_{\text{left}}$  and  $v_{\text{right}}$  using weight  $dx$ :

.12: Bilinear Weight Calculation

### **Problem 1.12: Bilinear Weight Calculation**

Write code to calculate bilinear interpolation weight for points  $p$  and  $q$ :

.13: Value Accumulation

### **Problem 1.13: Value Accumulation**

Write code to accumulate weighted interpolated value to  $\text{current\_value}$  using  $\text{weight}$ ,  $\text{mask}$ , and  $\text{interpolated}$ :

.14: Corner Position Calculation

### **Problem 1.14: Corner Position Calculation**

Write code to calculate all four corner positions  $(y_0, x_0)$ ,  $(y_0, x_1)$ ,  $(y_1, x_0)$ ,  $(y_1, x_1)$  from  $q_y$  and  $q_x$ :

.15: Mask Application

### **Problem 1.15: Mask Application**

Write code to apply modulation mask  $m_k$  to an interpolated value:

## TASK 2: CNN PYTORCH - WRITE CODE BLOCKS

.1: Conv Layer Definition

### Problem 2.1: Conv Layer Definition

Write code to define a Conv2d layer with 32 input channels, 64 output channels, kernel size 3, padding 1:

.2: BatchNorm Layer Definition

### Problem 2.2: BatchNorm Layer Definition

Write code to define a BatchNorm2d layer for 128 channels:

.3: MaxPool Layer Definition

### Problem 2.3: MaxPool Layer Definition

Write code to define a MaxPool2d layer that halves spatial dimensions:

.4: Linear Layer Definition

### Problem 2.4: Linear Layer Definition

Write code to define a Linear layer with 4096 inputs and 512 outputs:

.5: Dropout Layer Definition

### Problem 2.5: Dropout Layer Definition

Write code to define a Dropout layer with probability 0.5:

.6: ReLU Activation

### **Problem 2.6: ReLU Activation**

Write code to apply ReLU activation to variable x:

.7: Tensor Flattening

### **Problem 2.7: Tensor Flattening**

Write code to flatten tensor x while preserving batch dimension:

.8: Device Transfer

### **Problem 2.8: Device Transfer**

Write code to move tensor images to device:

.9: Optimizer Zero Grad

### **Problem 2.9: Optimizer Zero Grad**

Write code to clear gradients from optimizer:

.10: Forward Pass

### **Problem 2.10: Forward Pass**

Write code to perform forward pass through model with input images:

.11: Loss Calculation

### **Problem 2.11: Loss Calculation**

Write code to calculate loss using criterion with outputs and labels:

.12: Backward Pass

### **Problem 2.12: Backward Pass**

Write code to perform backward pass on loss:

.13: Optimizer Step

### **Problem 2.13: Optimizer Step**

Write code to update model parameters using optimizer:

.14: Top-1 Prediction

### **Problem 2.14: Top-1 Prediction**

Write code to get top-1 predictions from outputs:

.15: Accuracy Calculation

### **Problem 2.15: Accuracy Calculation**

Write code to calculate accuracy from predicted and labels tensors:

.16: Model Training Mode

### **Problem 2.16: Model Training Mode**

Write code to set model to training mode:

.17: Model Evaluation Mode

### **Problem 2.17: Model Evaluation Mode**

Write code to set model to evaluation mode:

.18: No Gradient Context

### **Problem 2.18: No Gradient Context**

Write code to create context where gradients are disabled:

.19: Dataset Size Calculation

### **Problem 2.19: Dataset Size Calculation**

Write code to calculate training size as 80% of total dataset size:

.20: Random Split

### **Problem 2.20: Random Split**

Write code to split dataset into `train_set` and `val_set` with sizes `train_size` and `val_size`:

## TASK 3: RNN - WRITE CODE BLOCKS

.1: Character Set Creation

### Problem 3.1: Character Set Creation

Write code to create sorted list of unique characters from text:

.2: Character to Index Mapping

### Problem 3.2: Character to Index Mapping

Write code to create dictionary mapping each character to its index:

.3: Index to Character Mapping

### Problem 3.3: Index to Character Mapping

Write code to create dictionary mapping each index to its character:

.4: Input Sequence Creation

### Problem 3.4: Input Sequence Creation

Write code to create input sequence (all characters except last) from text:



.5: Target Sequence Creation

### **Problem 3.5: Target Sequence Creation**

Write code to create target sequence (all characters except first) from text:

.6: One-Hot Vector Creation

### **Problem 3.6: One-Hot Vector Creation**

Write code to create one-hot vector of given size with 1 at given index:

.7: Input-to-Hidden Weight Initialization

### **Problem 3.7: Input-to-Hidden Weight Initialization**

Write code to initialize  $W_{xh}$  weight matrix for RNN with hidden size  $H$  and vocab size  $V$ :

.8: Hidden-to-Hidden Weight Initialization

### **Problem 3.8: Hidden-to-Hidden Weight Initialization**

Write code to initialize  $W_{hh}$  weight matrix for RNN with hidden size  $H$ :

.9: Hidden Bias Initialization

### Problem 3.9: Hidden Bias Initialization

Write code to initialize bias vector  $b_{xh}$  for hidden layer with size  $H$ :

.10: Output Weight Initialization

### Problem 3.10: Output Weight Initialization

Write code to initialize  $W_{hy}$  weight matrix from hidden to output with vocab size  $V$  and hidden size  $H$ :

.11: Hidden State Initialization

### Problem 3.11: Hidden State Initialization

Write code to initialize hidden state vector with zeros of size  $H$ :

.12: Input Contribution Calculation

### Problem 3.12: Input Contribution Calculation

Write code to calculate input contribution to hidden state using  $W_{xh}$  and  $x_t$ :

.13: Hidden Recurrence Calculation

### Problem 3.13: Hidden Recurrence Calculation

Write code to calculate recurrent contribution to hidden state using  $W_{hh}$  and previous hidden state  $h$ :

.14: Hidden State Update

### Problem 3.14: Hidden State Update

Write code to update hidden state using  $\tanh$  activation with all contributions and biases:

.15: Output Logits Calculation

### Problem 3.15: Output Logits Calculation

Write code to calculate output logits  $s_t$  from hidden state  $h$  using  $W_{hy}$  and  $b_y$ :

.16: Character Index Lookup

### Problem 3.16: Character Index Lookup

Write code to get index of character 'e' from `char2idx` dictionary:

.17: Input List Creation

### Problem 3.17: Input List Creation

Write code to convert input sequence to list of one-hot vectors:

.18: Target Tensor Creation

### **Problem 3.18: Target Tensor Creation**

Write code to convert target sequence to tensor of character indices:

.19: Logits Stacking

### **Problem 3.19: Logits Stacking**

Write code to stack list of logits tensors into single tensor:

.20: Log Softmax Calculation

### **Problem 3.20: Log Softmax Calculation**

Write code to calculate log softmax of logits along vocabulary dimension:

.21: NLL Loss Calculation

### **Problem 3.21: NLL Loss Calculation**

Write code to calculate negative log likelihood loss from log\_probs and targets:

.22: Gradient Calculation

### **Problem 3.22: Gradient Calculation**

Write code to calculate gradient of loss with respect to  $W_{xh}$ :

## DATA PREPROCESSING - WRITE CODE BLOCKS

.1: CIFAR-100 Mean Values

### Problem 4.1: CIFAR-100 Mean Values

Write code to define CIFAR-100 normalization mean values:

.2: CIFAR-100 Std Values

### Problem 4.2: CIFAR-100 Std Values

Write code to define CIFAR-100 normalization standard deviation values:

.3: ToTensor Transform

### Problem 4.3: ToTensor Transform

Write code to create ToTensor transform:

.4: Normalization Transform

### Problem 4.4: Normalization Transform

Write code to create Normalize transform with CIFAR-100 mean and std:

.5: Random Crop Transform

## Problem 4.5: Random Crop Transform

Write code to create RandomCrop transform with size 32 and padding 4:

.6: Random Flip Transform

## Problem 4.6: Random Flip Transform

Write code to create RandomHorizontalFlip transform:

.7: Transform Composition

## Problem 4.7: Transform Composition

Write code to compose multiple transforms into single transform:

.8: CIFAR-100 Dataset Loading

## Problem 4.8: CIFAR-100 Dataset Loading

Write code to load CIFAR-100 training dataset with transform:

.9: DataLoader Creation

## Problem 4.9: DataLoader Creation

Write code to create DataLoader with batch size 128 and shuffle=True:



## OPTIMIZATION AND TRAINING - WRITE CODE BLOCKS

.1: CrossEntropy Loss Definition

### Problem 5.1: CrossEntropy Loss Definition

Write code to define CrossEntropy loss function:

.2: SGD Optimizer Definition

### Problem 5.2: SGD Optimizer Definition

Write code to define SGD optimizer with learning rate 0.01 and momentum 0.9:

.3: Adam Optimizer Definition

### Problem 5.3: Adam Optimizer Definition

Write code to define Adam optimizer with learning rate 0.001:

.4: Model Parameter Count

### Problem 5.4: Model Parameter Count

Write code to count total number of parameters in model:

.5: Learning Rate Update

### **Problem 5.5: Learning Rate Update**

Write code to multiply learning rate by 0.1 for all parameter groups:

.6: Model State Save

### **Problem 5.6: Model State Save**

Write code to save model state dictionary to file 'model.pth':

.7: Model State Load

### **Problem 5.7: Model State Load**

Write code to load model state dictionary from file 'model.pth':

.8: Gradient Clipping

### **Problem 5.8: Gradient Clipping**

Write code to clip gradients to maximum norm of 1.0:

3cm

.9: Top-5 Accuracy

### **Problem 5.9: Top-5 Accuracy**

Write code to calculate top-5 accuracy from outputs and labels:

.10: Loss Item Extraction

## Problem 5.10: Loss Item Extraction

Write code to extract scalar value from loss tensor:

## DEBUGGING AND UTILITIES - WRITE CODE BLOCKS

.1: Tensor Shape Check

### Problem 6.1: Tensor Shape Check

Write code to print shape of tensor x:

.2: Tensor Device Check

### Problem 6.2: Tensor Device Check

Write code to check which device tensor x is on:

.3: Model Device Transfer

### Problem 6.3: Model Device Transfer

Write code to move entire model to GPU:

.4: Gradient Existence Check

### Problem 6.4: Gradient Existence Check

Write code to check if parameter has gradients:

.5: Memory Usage Check

### Problem 6.5: Memory Usage Check

Write code to check CUDA memory usage:

.6: Random Seed Setting

### **Problem 6.6: Random Seed Setting**

Write code to set PyTorch random seed to 42:

.7: Numpy Seed Setting

### **Problem 6.7: Numpy Seed Setting**

Write code to set numpy random seed to 42:

.8: Model Summary

### **Problem 6.8: Model Summary**

Write code to print model architecture:

.9: Batch Dimension Check

### **Problem 6.9: Batch Dimension Check**

Write code to get batch size from tensor x:

.10: Tensor Type Conversion

### **Problem 6.10: Tensor Type Conversion**

Write code to convert tensor x to float type:

## ANSWER BANK

Task 1 - Deformable CNN Answers:

### Problem Task 1 - Deformable CNN Answers:

```
1.1: y0 = int(np.floor(q_y))
1.2: x1 = int(np.ceil(q_x))
1.3: k = kh * K_w + kw
1.4: delta_y = delta[n, 2 * k, h_out, w_out]
1.5: delta_x = delta[n, 2 * k + 1, h_out, w_out]
1.6: h_start = h_out * stride
1.7: sample_y = h_start + kh * dilation + delta_y
1.8: if 0 <= y < H and 0 <= x < W:
1.9: value = img[y, x] if (0 <= y < H and 0 <= x < W) else 0.0
1.10: dx = q_x - x0
1.11: result = v_left * (1 - dx) + v_right * dx
1.12: weight = (1 - abs(p_x - q_x)) * (1 - abs(p_y - q_y))
1.13: current_value += weight * mask * interpolated
1.14: y0, x0 = int(np.floor(q_y)), int(np.floor(q_x))
    y1, x1 = y0 + 1, x0 + 1
1.15: modulated_value = m_k * interpolated_value
```

Task 2 - CNN PyTorch Answers:

### Problem Task 2 - CNN PyTorch Answers:

```
2.1: self.conv = nn.Conv2d(32, 64, kernel_size=3, padding=1)
2.2: self.bn = nn.BatchNorm2d(128)
2.3: self.pool = nn.MaxPool2d(2, 2)
2.4: self.fc = nn.Linear(4096, 512)
2.5: self.dropout = nn.Dropout(0.5)
2.6: x = F.relu(x)
2.7: x = x.view(x.size(0), -1)
2.8: images = images.to(device)
2.9: optimizer.zero_grad()
2.10: outputs = model(images)
2.11: loss = criterion(outputs, labels)
```

```

2.12: loss.backward()

2.13: optimizer.step()

2.14: _, predicted = torch.max(outputs, 1)

2.15: accuracy = (predicted == labels).sum().item() / labels.size(0) * 100

2.16: model.train()

2.17: model.eval()

2.18: with torch.no_grad():

2.19: train_size = int(0.8 * len(dataset))

2.20: train_set, val_set = random_split(dataset, [train_size, val_size])

```

Task 3 - RNN Answers:

## Problem Task 3 - RNN Answers:

```

3.1: chars = sorted(list(set(text)))

3.2: char2idx = {ch: i for i, ch in enumerate(chars)}

3.3: idx2char = {i: ch for i, ch in enumerate(chars)}

3.4: input_seq = text[:-1]

3.5: target_seq = text[1:]

3.6: vec = torch.zeros(size)
vec[idx] = 1.0

3.7: W_xh = torch.randn(H, V, requires_grad=True) * 0.1

3.8: W_hh = torch.randn(H, H, requires_grad=True) * 0.1

3.9: b_xh = torch.zeros(H, requires_grad=True)

3.10: W_hy = torch.randn(V, H, requires_grad=True) * 0.1

3.11: h = torch.zeros(H)

3.12: input_contrib = W_xh @ x_t

3.13: hidden_contrib = W_hh @ h

3.14: h = torch.tanh(W_xh @ x_t + b_xh + W_hh @ h + b_hh)

3.15: s_t = W_hy @ h + b_y

3.16: idx = char2idx['e']

3.17: inputs = [one_hot(char2idx[ch], V) for ch in input_seq]

3.18: targets = torch.tensor([char2idx[ch] for ch in target_seq], dtype=torch.long)

3.19: logits = torch.stack(logits_list)

3.20: log_probs = F.log_softmax(logits, dim=1)

```



```
3.21: loss = F.nll_loss(log_probs, targets)
```

```
3.22: grad.W_xh = torch.autograd.grad(loss, W_xh, retain_graph=True)[0]
```

Additional Sections Answers:

## Problem Additional Sections Answers:

```
4.1: mean = [0.5071, 0.4867, 0.4408]
```

```
4.2: std = [0.2675, 0.2565, 0.2761]
```

```
4.3: transform = transforms.ToTensor()
```

```
4.4: normalize = transforms.Normalize(mean=[0.5071, 0.4867, 0.4408], std=[0.2675, 0.2565, 0.2761])
```

```
4.5: crop = transforms.RandomCrop(32, padding=4)
```

```
4.6: flip = transforms.RandomHorizontalFlip()
```

```
4.7: transform = transforms.Compose([transform1, transform2, ...])
```

```
4.8: dataset = CIFAR100(root='./data', train=True, transform=transform)
```

```
4.9: loader = DataLoader(dataset, batch_size=128, shuffle=True)
```

```
5.1: criterion = nn.CrossEntropyLoss()
```

```
5.2: optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9)
```

```
5.3: optimizer = optim.Adam(model.parameters(), lr=0.001)
```

```
5.4: total_params = sum(p.numel() for p in model.parameters())
```

```
5.5: for param_group in optimizer.param_groups: param_group['lr'] *= 0.1
```

```
5.6: torch.save(model.state_dict(), 'model.pth')
```

```
5.7: model.load_state_dict(torch.load('model.pth'))
```

```
5.8: torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)
```

```
5.9: _, top5_pred = outputs.topk(5, dim=1)
top5_acc = top5_pred.eq(labels.view(-1, 1).expand_as(top5_pred)).sum().item()
```

```
5.10: loss_value = loss.item()
```

```
6.1: print(x.shape)
```

```
6.2: print(x.device)
```

```
6.3: model = model.to('cuda')
```

```
6.4: if param.grad is not None:
```

```
6.5: print(torch.cuda.memory_allocated())
```

```
6.6: torch.manual_seed(42)
```

```
6.7: np.random.seed(42)
```

```
6.8: print(model)
```

```
6.9: batch_size = x.size(0)
```

```
6.10: x = x.float()
```

Additional Answers

## Problem Additional Answers

C1: 32 x 32, C2: 8 x 8, C3: 4096, C4: 73856, C5: 256

S1: numpy, torch.nn, F, transforms

S2: nn.Module, super

S3: def forward(self, x)

S4: nn.CrossEntropyLoss()

S5: optim.SGD(model.parameters(), lr=0.01, momentum=0.9)

M1: x = self.dropout(x) if self.training else x

M2: for param\_group in optimizer.param\_groups: param\_group['lr'] \*= 0.1

M3: torch.save(model.state\_dict(), 'model.pth')

M4: outputs = model(images.to(device))

M5: correct = (predicted == labels).sum().item()